


Concurrency Control

Introduction

- Interactions among the transactions can cause the database state inconsistent.
- Each transaction
 - Individually ensures consistent state with no system failure.
- When several transactions are processed in parallel, inconsistency may occur.
 - Needs regulation
 - Scheduler controls the accesses
 - Scheduler is a protocol
 - Question: how to design a protocol ?

Scheduler



- Transaction passes reads and write requests to scheduler
- Scheduler executes directly, if the elements are present in the buffer.
 - Otherwise, data is brought to memory.
- Scheduler may delay or abort the transaction.
 - What is the criteria ?
 - Read Uncommitted, Read committed, Repeatable read, Serializability.
 - What is the protocol ?
 - No locking, Locking, optimistic, timestamp, and so on

Dirty read, non-repeatable read and phantom problems

- Dirty read problem: No locks: dirty read problem occurs
 - A dirty read occurs when a transaction is allowed to read data from a row that has been modified by another running transaction and not yet committed.
- Non repeatable read problem: No read locks.
 - A *non-repeatable read* occurs, when during the course of a transaction, a row is retrieved twice and the values within the row differ between reads.
- Phantom problem: read locks are released as soon as possible
 - Execution two similar queries give different results.

Isolation levels and read phenomena

Isolation level	Dirty reads	Non-repeatable reads	Phantoms
Read Uncommitted	X	X	X
Read Committed	-	X	X
Repeatable Read	-	-	X
Serializable	-	-	-

Note: "X" means that the isolation level suffers that phenomenon, while "-" means that it does not suffer it.

Refer: [http://en.wikipedia.org/wiki/Isolation_\(database_systems\)](http://en.wikipedia.org/wiki/Isolation_(database_systems))

Outline

- Serial and serializable schedules
- Conflict serializability
- Enforcing serializability by locks
- Locking system with several lock modes
- An architecture of a locking scheduler
- Managing hierarchies of database elements
- The tree protocol
- Concurrency control by timestamps
- Concurrency control by validation

Serial and Serializable Schedules.

- Correctness principle.
 - If executed alone, it starts in a consistent state and leaves the database in a consistent state, when the transaction ends.
- What is the correctness principle if several transactions are executed concurrently ?
- What are the protocols which are to be followed by scheduler to ensure the correctness ?

Schedules (or History)

- A schedule is a time-ordered sequence of the important actions taken by one or more transactions.
 - We consider read and write operations that occur in the buffer.
 - Ignore INPUT and OUTPUT operations

Serial Schedules

- A schedule is serial, if its actions consist of all the actions of one transaction, another transaction and so on.
 - A schedule S is serial if any two transactions T , and T' , if any action of T precedes any action of T' , then all the actions of T precede all actions of T' .

Example:

T1: Read(A)
 $A \leftarrow A+100$
 Write(A)
 Read(B)
 $B \leftarrow B+100$
 Write(B)

T2: Read(A)
 $A \leftarrow A \times 2$
 Write(A)
 Read(B)
 $B \leftarrow B \times 2$
 Write(B)

Constraint: $A=B$

Schedule A

T1

T2

Read(A); $A \leftarrow A+100$

Write(A);

Read(B); $B \leftarrow B+100$;

Write(B);

Read(A); $A \leftarrow A \times 2$;

Write(A);

Read(B); $B \leftarrow B \times 2$;

Write(B);

Schedule A

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A+100$			
Write(A);		125	
Read(B); $B \leftarrow B+100$;			
Write(B);		125	
	Read(A); $A \leftarrow A \times 2$;		
	Write(A);	250	
	Read(B); $B \leftarrow B \times 2$;		
	Write(B);	250	
		250	250

Serial Schedule in which T1 precedes T2

Schedule B

T1

T2

Read(A); $A \leftarrow A \times 2$;

Write(A);

Read(B); $B \leftarrow B \times 2$;

Write(B);

Read(A); $A \leftarrow A + 100$

Write(A);

Read(B); $B \leftarrow B + 100$;

Write(B);

Another serial schedule: T2 precedes T1

Schedule B

T1	T2	A	B
		25	25
	Read(A); $A \leftarrow A \times 2$;		
	Write(A);	50	
	Read(B); $B \leftarrow B \times 2$;		
	Write(B);	50	
Read(A); $A \leftarrow A + 100$			
Write(A);		150	
Read(B); $B \leftarrow B + 100$;			
Write(B);		150	
		150	150

Serializable Schedules

- Correctness principle
 - Every serial schedule ensures the correctness.
- Are there other schedules which preserve the consistency ?
 - These are called serializable schedules which are equivalent to a serial schedule.

Schedule C

T1

T2

Read(A); $A \leftarrow A+100$

Write(A);

Read(B); $B \leftarrow B+100$;

Write(B);

Read(A); $A \leftarrow A \times 2$;

Write(A);

Read(B); $B \leftarrow B \times 2$;

Write(B);

Schedule C

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A+100$			
Write(A);		125	
	Read(A); $A \leftarrow A \times 2$;		
	Write(A);	250	
Read(B); $B \leftarrow B+100$;			
Write(B);		125	
	Read(B); $B \leftarrow B \times 2$;		
	Write(B);	250	
		250	250

Serializable, but not serial

Schedule D

T1	T2
Read(A); $A \leftarrow A+100$ Write(A);	Read(A); $A \leftarrow A \times 2$; Write(A);
Read(B); $B \leftarrow B+100$; Write(B);	Read(B); $B \leftarrow B \times 2$; Write(B);

Not serial not serializable

Schedule D

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A+100$		125	
Write(A);		250	
	Read(A); $A \leftarrow A \times 2$;	50	
	Write(A);	150	
	Read(B); $B \leftarrow B \times 2$;	250	150
	Write(B);		
Read(B); $B \leftarrow B+100$;			
Write(B);			

The effect of transaction semantics

- The details of transactions matter (see the next example).
 - But it is very difficult to analyze the semantics.

Schedule E

Same as Schedule D
but with new T2'

T1	T2'
Read(A); $A \leftarrow A+100$ Write(A);	
	Read(A); $A \leftarrow A \times 1$; Write(A);
	Read(B); $B \leftarrow B \times 1$; Write(B);
Read(B); $B \leftarrow B+100$; Write(B);	

A schedule is serializable due to behavior of transactions

Schedule E

Same as Schedule D
but with new T2'

		A	B
T1	T2'	25	25
Read(A); $A \leftarrow A+100$			
Write(A);		125	
	Read(A); $A \leftarrow A \times 1$;		
	Write(A);	125	
	Read(B); $B \leftarrow B \times 1$;		
	Write(B);	25	
Read(B); $B \leftarrow B+100$;			
Write(B);		125	
		125	125

Notations

- $r_T(X)$: transaction T reads database element X
- $w_T(X)$: transaction T writes database element X
- $r_i(X)$ is same as $r_{T_i}(X)$:
- $w_i(X)$ is same as $w_{T_i}(X)$:
- Example

$T1 = r_1(A); w_1(A); r_1(B); w_1(B);$

$T2 = r_2(A); w_2(A); r_2(B); w_2(B);$

- Want schedules , regardless of
 - initial state and
 - transaction semantics
- Only look at order of read and writes

Example: Consider serializable schedule


$S_c = r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B)$

Notation for Transactions and Schedules

- An action is $ri(X)$ or $wi(X)$
- T_i is a sequence of actions with subscript “i”.
- A schedule S of a set of transactions T is a sequence of actions, in which for each T_i in T , the actions of T_i appear in S in the **same order** as in the definition of “ T_i ” itself.
- Example:
$$Sc = r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B);$$

Example:

$$Sc = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$$



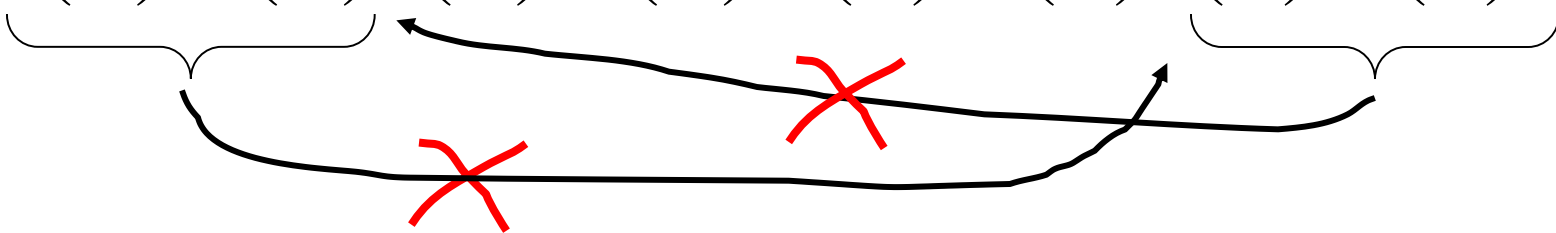
$$Sc' = r_1(A)w_1(A) \underbrace{r_1(B)w_1(B)}_{T_1} \underbrace{r_2(A)w_2(A)r_2(B)w_2(B)}_{T_2}$$

T_1

T_2

However, for S_d :

$S_d = r_1(A)w_1(A)r_2(A)w_2(A) \ r_2(B)w_2(B)r_1(B)w_1(B)$



- as a matter of fact,
 T_2 must precede T_1
 in any equivalent schedule,
 i.e., $T_2 \rightarrow T_1$

Conflict Serializability

- Scheduler ensures that schedules are serializable by ensuring a condition conflict-serializability.
- It is based on the idea of conflict
- Conflict: Let T_i and T_j are transactions
 - $ri(X); rj(Y)$ is never conflict even $X=Y$ as they do not change any value
 - $ri(X); wj(Y)$ is not a conflict provided that $X \neq Y$.
 - $wi(X); rj(Y)$ is not a conflict if $X \neq Y$.
 - $wi(X); wj(Y)$ is not a conflict if $X \neq Y$
- The actions of the same transaction conflict.; $ri(X); wi(Y)$ conflict; order is fixed by a transaction. DBMS can not reorder!
- Two writes of different transaction on the same element conflict: $wi(X), wj(X)$ conflict.
- Read and write of the same database element by different transactions also conflict. $ri(X); wj(X)$ conflict. Also, $wi(X)$ and $ri(X)$ conflict.

Conflict Serializability

- Two actions of different transactions may be swapped unless
 - They involve the same database element, and
 - at least one of them is write.
- Carry out the non conflicting swaps and try to convert the schedule into a serial schedule.
- Conflict-equivalent
 - Two schedules are conflict equivalent if they can be turned one into other by a sequence of non-conflicting swaps of **ADJACENT ACTIONS**.
- Conflict serializable
 - If a schedule is conflict equivalent to a serial schedule.
- Conflict serializable schedule is a serializable schedule.
- Note:
 - Conflict serializability is not required for a schedule to be serializable.
 - But many commercial system use serializability criteria to guarantee serializabilty

Precedence Graphs and a test for conflict serializability

- Nodes: transactions in S

Arcs: $T_i \rightarrow T_j$ whenever

- $p_i(A), q_j(A)$ are actions in S
- $p_i(A) <_S q_j(A)$
- at least one of p_i, q_j is a write

- Insert the edge
 - $ri(X), wj(X)$
 - $wj(X), ri(X)$
 - $wi(X), wj(X)$
 - $wj(X), wi(X)$

Exercise:

- What is $P(S)$ for
 $S = w_3(A) \ w_2(C) \ r_1(A) \ w_1(B) \ r_1(C) \ w_2(A) \ r_4(A) \ w_4(D)$
- Is S serializable?

Another Exercise:

- What is $P(S)$ for
 $S = w_1(A) \ r_2(A) \ r_3(A) \ w_4(A) \ ?$

Enforcing serializability by locks

- How to enforce serializable schedules?

Option 1: run system, recording


P(S);	at end of day, check for
-------	--------------------------

P(S)	cycles and declare if
------	-----------------------

execution	was good
-----------	----------

How to enforce serializable schedules?

Option 2: prevent P(S) cycles from occurring



Serializability with Locks

- Arbitrary execution may result into non-serializable schedule.
- Job of scheduler is to prevent orders of actions that lead to an non-serializable schedule.
- We first discuss the lock-based scheduler with single lock
 - Lock must be obtained for a database element before accessing it.

Proper use of locks: Legal schedule

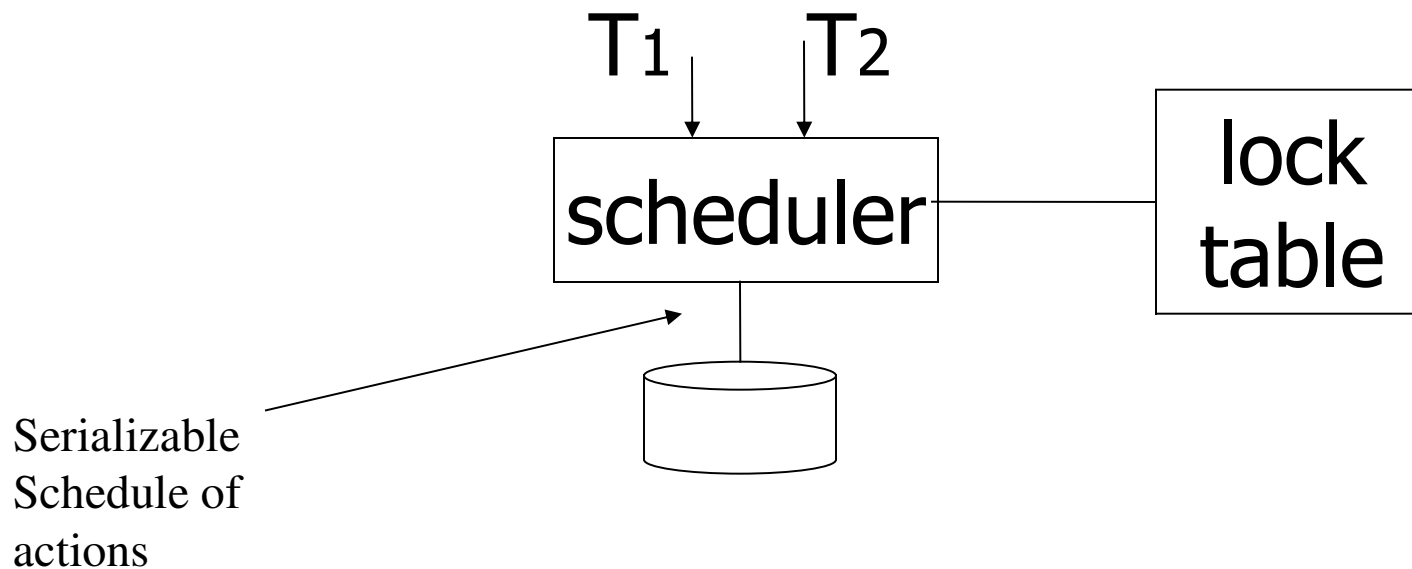
- Consistency of transactions
 - Transaction can only read/write an element if it has previously has requested the lock.
 - If transaction locks an element, it must unlock that element.
- No two transactions must not lock the given object simultaneously. Lock can be given after the first transaction unlocking it.

A locking protocol (with one kind of lock)

Two new actions:

lock (exclusive): $li(A)$ T_i requests a lock

unlock: $ui(A)$ T_i releases the lock



Consistency Condition for transactions


- If there are actions $li(X)$ followed by $lj(X)$ in a schedule, there should be $ui(X)$ action between these actions.

Rule #1: Well-formed transactions: transaction should request a lock and release the lock

$T_i: \dots li(A) \dots pi(A) \dots ui(A) \dots$

Rule #2 Legal scheduler: No two transactions should not lock the element without the first unlocked it.

$S = \dots\dots\dots l_i(A) \dots\dots\dots u_i(A) \dots\dots\dots$



no $l_j(A)$

Exercise:

- What schedules are legal?

What transactions are well-formed?

$S1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

$S2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$

$S3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$
 $l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

Exercise:

- What schedules are legal?

What transactions are well-formed?

$S1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

$S2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$

$S3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$
 $l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

Schedule F

T1	T2
$l_1(A); \text{Read}(A)$ $A \leftarrow A + 100; \text{Write}(A); u_1(A)$	$l_2(A); \text{Read}(A)$ $A \leftarrow A \times 2; \text{Write}(A); u_2(A)$ $l_2(B); \text{Read}(B)$ $B \leftarrow B \times 2; \text{Write}(B); u_2(B)$
$l_1(B); \text{Read}(B)$ $B \leftarrow B + 100; \text{Write}(B); u_1(B)$	

Schedule F

		A	B
T1	T2	25	25
l ₁ (A);Read(A)			
A ← A+100;Write(A);u ₁ (A)		125	
	l ₂ (A);Read(A)		
	A ← Ax2;Write(A);u ₂ (A)	250	
	l ₂ (B);Read(B)		
	B ← Bx2;Write(B);u ₂ (B)		50
l ₁ (B);Read(B)			
B ← B+100;Write(B);u ₁ (B)			150
		250	150

Schedule F

		A	B
T1	T2	25	25
$l_1(A); \text{Read}(A)$ $A \leftarrow A + 100; \text{Write}(A); u_1(A)$		125	
	$l_2(A); \text{Read}(A)$ $A \leftarrow A \times 2; \text{Write}(A); u_2(A)$	250	
	$l_2(B); \text{Read}(B)$ $B \leftarrow B \times 2; \text{Write}(B); u_2(B)$		50
			150
$l_1(B); \text{Read}(B)$ $B \leftarrow B + 100; \text{Write}(B); u_1(B)$		250	150

Note: Schedule F is a legal schedule of consistent transactions, but it is not serializable.

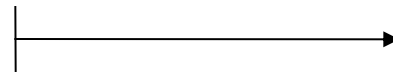
Two-phase locking


- Rule #3: In every transaction, all lock requests precede all unlock requests.

Rule #3 Two phase locking (2PL) for transactions

$T_i = \dots \dots \text{li}(A) \dots \dots \text{ui}(A) \dots \dots$


no unlocks


no locks



Schedule G


T1	T2
$l_1(A); \text{Read}(A)$	
$A \leftarrow A + 100; \text{Write}(A)$	
$l_1(B); u_1(A)$	
	$l_2(A); \text{Read}(A)$
	$A \leftarrow A \times 2; \text{Write}(A); l_2(B)$

delayed

Schedule G

T1	T2
$l_1(A); \text{Read}(A)$	
$A \leftarrow A + 100; \text{Write}(A)$	
$l_1(B); u_1(A)$	
	$l_2(A); \text{Read}(A)$
	$A \leftarrow A \times 2; \text{Write}(A); l_2(B)$
$\text{Read}(B); B \leftarrow B + 100$	
$\text{Write}(B); u_1(B)$	

delayed



Schedule G

T1

$l_1(A); \text{Read}(A)$

$A \leftarrow A + 100; \text{Write}(A)$

$l_1(B); u_1(A)$

$\text{Read}(B); B \leftarrow B + 100$

$\text{Write}(B); u_1(B)$

T2

$l_2(A); \text{Read}(A)$


$A \leftarrow A \times 2; \text{Write}(A); l_2(B)$

delayed

$l_2(B); u_2(A); \text{Read}(B)$

$B \leftarrow B \times 2; \text{Write}(B); u_2(B);$


Schedule H (T2 reversed)



Deadlock !

- Assume deadlocked transactions are rolled back
 - They have no effect
 - They do not appear in schedule

E.g., Schedule H =


This space intentionally
left blank!

Next step:

Show that rules #1,2,3 \Rightarrow conflict-
serializable
schedules

Conflict rules for $l_i(A)$, $u_i(A)$:

- $l_i(A)$, $l_j(A)$ conflict
- $l_i(A)$, $u_j(A)$ conflict

Note: no conflict $\langle u_i(A), u_j(A) \rangle$, $\langle l_i(A), r_j(A) \rangle$,...

Theorem Rules #1,2,3 \Rightarrow conflict
(2PL) serializable
schedule

Theorem Rules #1,2,3 \Rightarrow conflict
(2PL) serializable
schedule

To help in proof:

Definition $\text{Shrink}(Ti) = \text{SH}(Ti)$
= first unlock
action of Ti

Lemma

$$Ti \rightarrow Tj \text{ in } S \Rightarrow SH(Ti) <_S SH(Tj)$$

Lemma

$$T_i \rightarrow T_j \text{ in } S \Rightarrow SH(T_i) <_S SH(T_j)$$

Proof of lemma:

$T_i \rightarrow T_j$ means that

$$S = \dots p_i(A) \dots q_j(A) \dots; \quad p, q \text{ conflict}$$

By rules 1,2:

$$S = \dots p_i(A) \dots u_i(A) \dots l_j(A) \dots q_j(A) \dots$$

Lemma


$$T_i \rightarrow T_j \text{ in } S \Rightarrow SH(T_i) <_S SH(T_j)$$

Proof of lemma:

$T_i \rightarrow T_j$ means that

$$S = \dots p_i(A) \dots q_j(A) \dots; \quad p, q \text{ conflict}$$

By rules 1,2:

$$S = \dots p_i(A) \dots u_i(A) \dots l_j(A) \dots q_j(A) \dots$$


$$\text{By rule 3:} \quad SH(T_i) \qquad SH(T_j)$$

$$\text{So, } SH(T_i) <_S SH(T_j)$$

Theorem Rules #1,2,3 \Rightarrow conflict
(2PL) serializable
schedule

Proof:

(1) Assume $P(S)$ has cycle


$$T_1 \rightarrow T_2 \rightarrow \dots T_n \rightarrow T_1$$

(2) By lemma: $SH(T_1) < SH(T_2) < \dots < SH(T_1)$

(3) Impossible, so $P(S)$ acyclic

(4) $\Rightarrow S$ is conflict serializable

2PL subset of Serializable



Locking Systems with Several Lock Modes

- In one lock scheme even “reading” action also requires a lock.
 - Several transactions can read X in parallel.
- We can use two kinds of locks
 - Shared lock or read lock
 - Exclusive lock or write lock
- sli(X): requests a shared lock on X
- xli(X): requests an exclusive lock on X
- ui(X): relinquishes the locks.

Requirements

- 1. For writing, exclusive lock is required and for reading any lock is OK.
 - $ri(X)$ must be preceded by $sli(X)$ or $xli(X)$ with no intervening $ui(X)$.
 - $wi(X)$ must be preceded by $xli(X)$ with no intervening $ui(X)$.
- 2: Two phase rule: locking must precede unlocking
 - No action $sli(X)$ or $xli(X)$ can be preceded by an action $ui(X)$
- 3. An object can be locked by several on a shared mode or exclusively by one transaction.
 - If $xli(X)$ appears in a schedule, there will be no $xlj(X)$ or $slj(X)$ for some j other than i , without an intervening $ui(X)$.
 - If $sli(X)$ appears in a schedule, then there can not be following $xlj(X)$ without intervening $ui(X)$.

Compatible matrix

		Lock Requested	
		S	X
Lock Held	S	true	false
	X	false	false

Example

T1

T2

sl1(A);r1(A)

sl2(A);r2(A)

sl2(B); r2(B)

sl1(B); r1(B);

xl1(B) denied

u2(A);u2(B)

xl1(B); w1(B);

u1(A); u2(B)

Note: If T1 would have requested exclusive lock initially, it would have been rejected.

Upgrading Locks

- If T_i has a shared lock on X can upgrade to exclusive lock.

Upgrading locks: deadlock

T1

T2

s11(A)

s12(A)

x11(A) denied

x12(A) denied

More concurrency Update locks

- To avoid deadlock problem
- Update lock can only read and not to write
- Only update lock will be upgraded to write lock later.
- We can grant update lock even though transactions have shared lock on X but, Once we have an update lock, other locks (shared, exclusive, update) are denied.

Update Locks: Compatibility Matrix

	S	X	U
S	Yes	No	Yes
X	No	No	No
U	No	No	No

Upgrading locks: No deadlock Problem

T1

T2

ul1(A); r1(A);

xl1(A), w1(A); u1(A)

ul2(A) denied


ul2(A); r2(A);
xl2(A); w2(A); u2(A);

Increment Locks

- For operations that commute each other; two transactions add constants to each other; it does not matter which goes first.

Example: Increment lock

- Atomic increment action: $\text{INi}(A)$
 $\{\text{Read}(A); A \leftarrow A+k; \text{Write}(A)\}$
- $\text{INi}(A), \text{INj}(A)$ do not conflict!



Increment Locks: Compatibility Matrix

	S	X	I
S	Yes	No	No
X	No	No	No
I	No	No	Yes

Increment locks:

T1

T2

sl1(A); r1(A);

sl2(A);r2(A);
il2(B);inc2(B);

il1(B), inc1(B);

u2(A); u2(B);


u1(A); u1(B);

Architecture of a Lock Scheduler

(1) Don't trust transactions

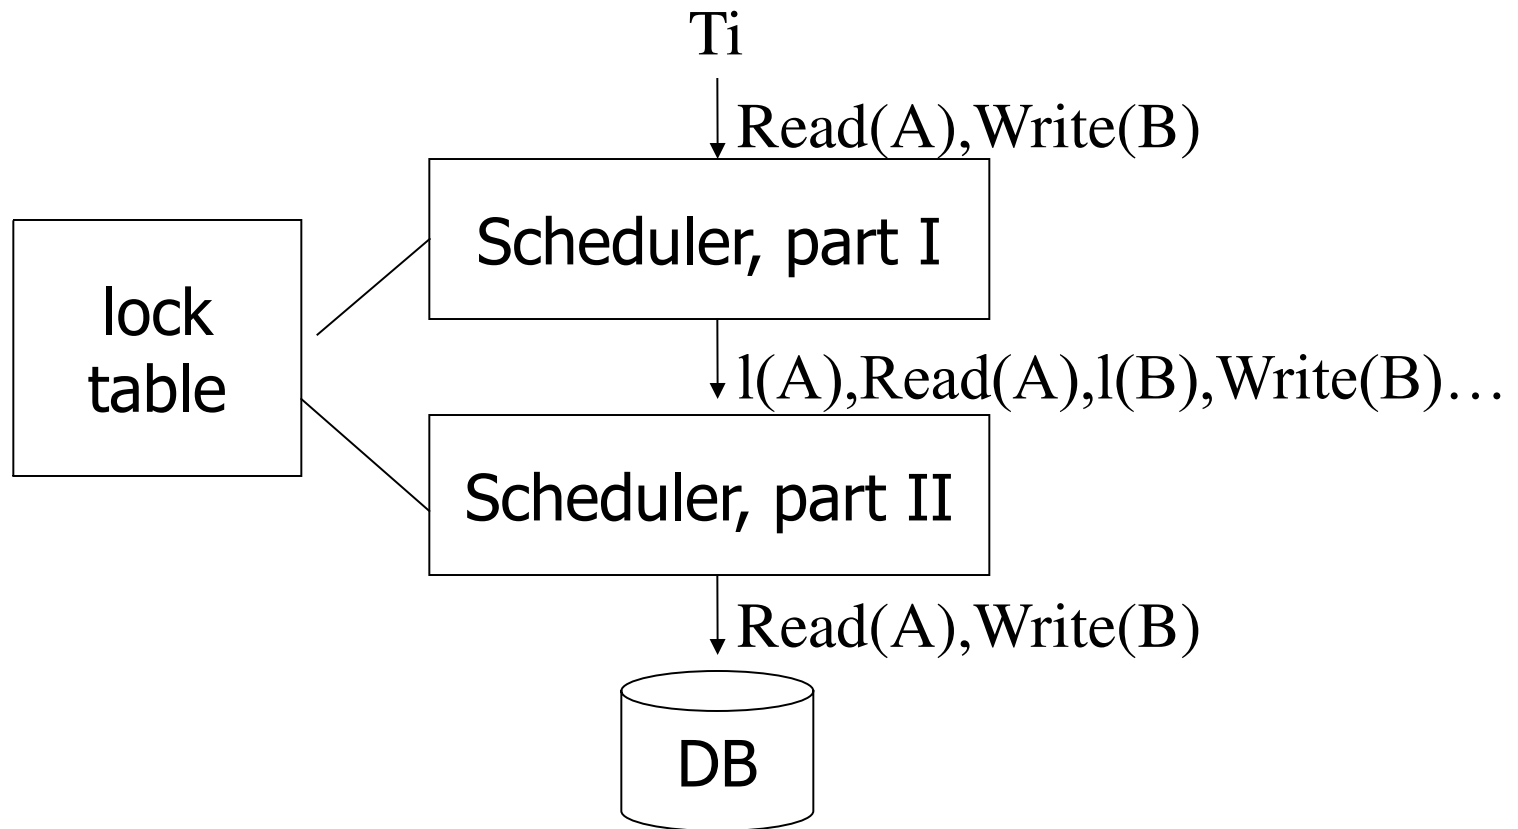
to request/release
locks

(2) Hold all locks until transaction
commits



Principles

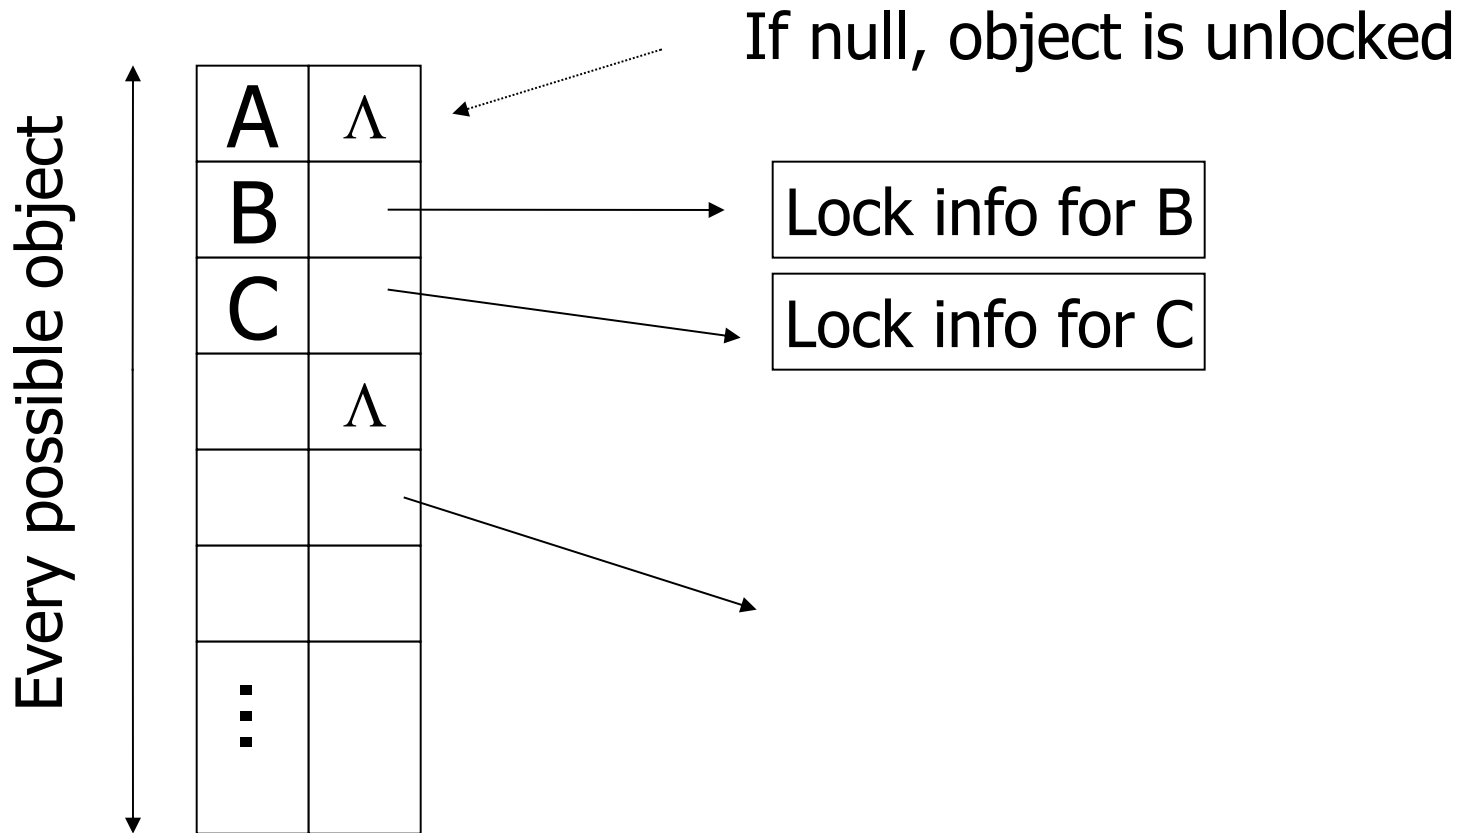
- Transactions do not request locks
 - Lock scheduler should insert the locks
- Transactions do not release the locks, the scheduler releases the locks based on commit or abort command.




Scheduler

- Part I selects appropriate mode of lock requests
- Part II executes the operations.
- When part I receives commit/abort by transaction manager, it releases the lock held by T. When a transaction is waiting for a lock, part I notifies part II.
- Part II starts executing waiting transactions.

Lock table Conceptually

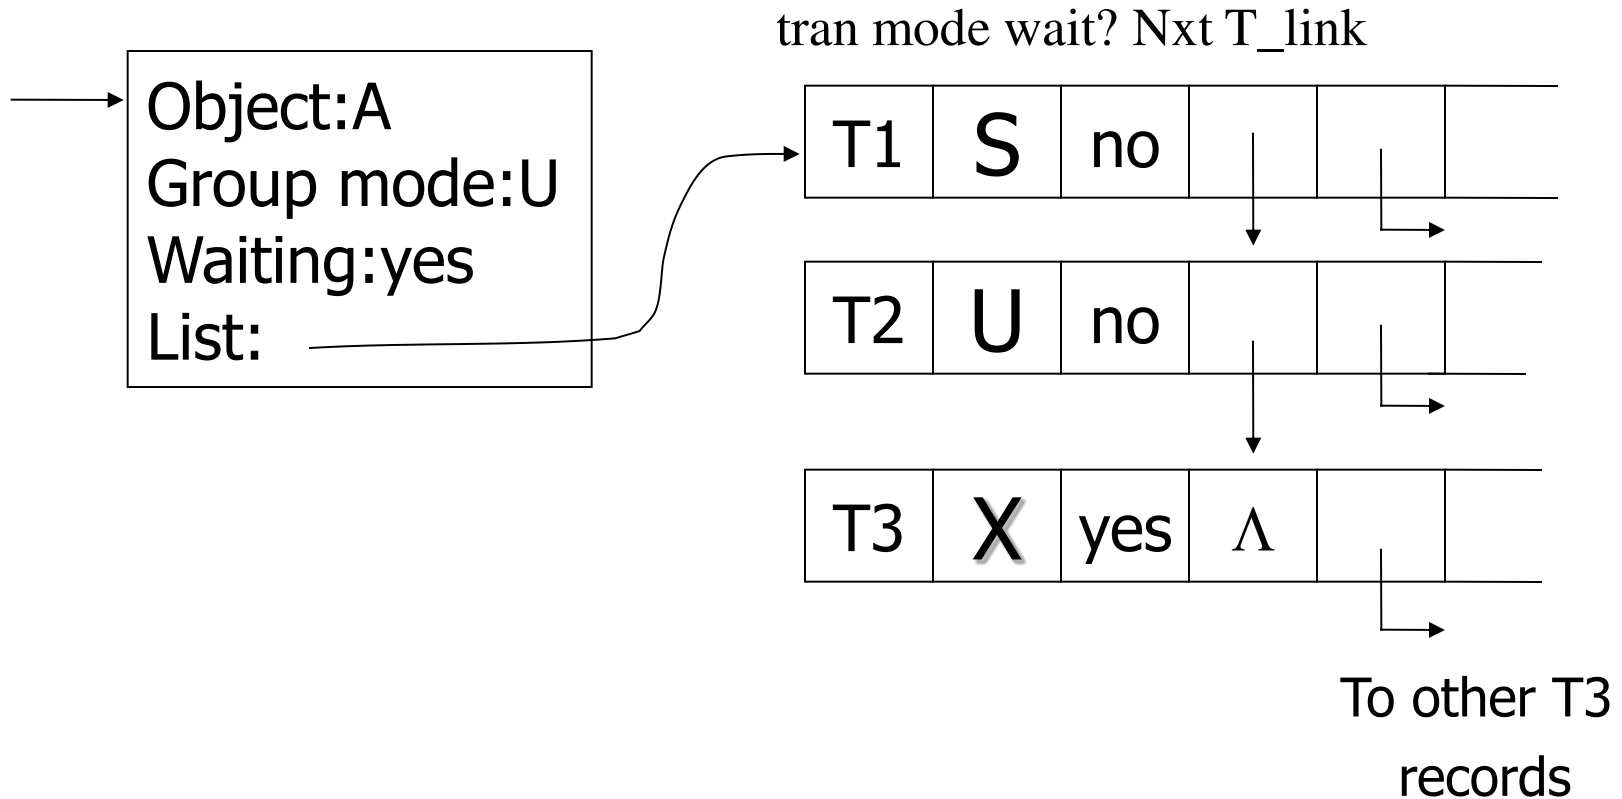


But use hash table:



If object not found in hash table, it is unlocked

Lock info for A - example



Selecting the list of requested locks

- FCFS
- Priority to shared locks
- Priority to upgrading

Managing hierarchies of database elements

- Two kinds of hierarchies
 - Hierarchy of lockable elements
- B-tree indexes

What are the objects we lock?

Relation A
Relation B
⋮

DB

Tuple A
Tuple B
Tuple C
⋮

DB

Disk block A
Disk block B
⋮

DB

?

- Locking works in any case, but should we choose small or large objects?

Locks with multiple granularity

- What is database element ?
 - One lock for each relation
 - One lock for each tuple
 - One lock for each page

Warning locks

Hierarchy Locking Protocol

- Relations are largest lockable elements
- Each relation consists of blocks
- Each block consists of tuples.
- The rules for managing hierarchy locking protocol involves warning locks and ordinary locks.

Warning protocol

1. To place an ordinary S or X lock on any element begin from the root
2. If we are at the element request S or X.
3. If the element is further down in the hierarchy, place a warning on this node.
 - If we want to request S lock further down, place IS lock on this node.
 - Here, “IS” means intention to obtain shared lock on the sub-element.
 - If we want to request X lock further down, place IX lock on this node. After granting the lock, proceed further and repeat 1,2, or 3.

Warning protocol: Compatibility matrix

	IS	IX	S	X
IS	Yes	Yes	Yes	No
IX	Yes	Yes	No	No
S	Yes	No	Yes	No
X	No	No	No	NO

Phantoms and Handling insertions Correctly


- We can only lock the existing elements!
- Find all Disney movies.
- But new Disney movie is inserted, in between.
Does not require the lock.
- Solution: lock the relation.

The Tree Protocol

- So far we have discussed about the elements organized in an hierarchical order
- Linked pattern of trees
 - Btrees
- Btrees allows the locking of individual nodes.
 - Treating entire B-tree as one element reduces the concurrency.
 - If we use S,X,Update, we can not design a CC protocol for Btree.
- If the transaction moves to child, and observes that split does not propagate upward, we can release the locks of higher level nodes.
 - Violates 2PL, but ensures serializability.


Example

- all objects accessed through root, following pointers




Example

- all objects accessed through root, following pointers





Example

- all objects accessed through root, following pointers





- can we release A lock if we no longer need A??



Idea: traverse like “Monkey Bars”



Idea: traverse like “Monkey Bars”




Idea: traverse like “Monkey Bars”



Why does this work?

- Assume all T_i start at root; exclusive lock
- $T_i \rightarrow T_j \Rightarrow T_i$ locks root before T_j




- Actually works if we don't always start at root

Rules: tree protocol (exclusive locks)

- (1) First lock by T_i may be on any item
- (2) After that, item Q can be locked by T_i
only if $\text{parent}(Q)$ locked by T_i
- (3) Items may be unlocked at any time
- (4) After T_i unlocks Q , it cannot relock Q ,
even it has a lock on node's parent.

- Tree-like protocols are used typically for B-tree concurrency control



E.g., during insert, do not release parent lock, until you are certain child does not have to split

Concurrency Control By Timestamps

Optimistic protocols

- Two methods
 - Timestamping
 - Serial schedule is according to timestamps.
 - Validation
 - Serial schedule is according to validation times.
- These are Optimistic protocols
 - Assume that conflicts are rare
- Locking protocols are pessimistic protocols
 - Assume that conflicts are frequent.

Timestamps

- It is a unique number
 - Transaction which starts later has higher timestamp.
- Each object X has two time stamps
 - $RT(X)$: read time of X : highest timestamp of the transaction which has read X .
 - $WT(X)$: write time of X : highest timestamp of the transaction which has written X .
 - $C(X)$, commit bit of X , true only most recent transaction to write X has already committed.
 - Suppose U writes X and aborts; dirty read problem can be avoided.

Timestamp-based Scheduling: Rules

- Suppose the scheduler receives $r_T(X)$
 - If $TS(T) \geq WT(X)$, read is possible
 - If $C(X)$ is true, grant the request
 - If $TS(T) > RT(X)$, set $RT(X) := TS(T)$; otherwise do not change $RT(X)$.
 - If $C(X)$ is false, delay T until $C(X)$ becomes true or the transaction which wrote X aborts
 - If $TS(T) < WT(X)$ rollback T and retransmit with new time stamp
- Suppose the scheduler receives $w_T(X)$
 - If $TS(T) \geq RT(X)$ and $TS(T) \geq WT(X)$, write is possible
 - Write X
 - Set $WT(X) := TS(T)$
 - Set $C(X)$ is false.
 - If $TS(T) \geq RT(X)$ and $TS(T) < WT(X)$, write is possible, but there is a later value of X . If $C(X)$ is true, ignore T . Otherwise delay T .
 - If $TS(T) < RT(X)$ then write is not possible
- If the scheduler receives $Commit(T)$, make $C(X)$ true.
- If the scheduler receives $Abort(T)$, take appropriate action.

Multiversion Timestamping

- Store the old versions
- Allow reads $r_T(X)$ that otherwise cause transaction T to abort

Multiversion Protocol

- When a new $w_T(X)$ occurs, new version of X is created. Its write time is $TS(T)$, let it be X_t .
- When a $rT(X)$ occurs, the scheduler finds version X_t of X such that $t \leq TS(T)$
- Write times are associated with the versions and they do not change
- Read times are associated with versions. Certain writes are rejected one whose time is less than the read time of the previous version.
- If there is no active transaction less than t , delete any version of X previous to X_t .

Timestamps and locking

- Timestamp is better if most of the transactions are read-only.
- Locking is better in high conflict situations.
- Approach followed by commercial systems
 - Read-only transactions are executed using multiversion timestamping.
 - Read/write transactions are executed using 2PL.

Concurrency Control by validation

- It is another type of optimistic control.
- Different from timestamping
 - Scheduler maintains the record of active transactions (not the timestamps of elements)
 - Before writing, a transaction goes through validation phase.
 - The items it wants to write are compared with write sets of other active transactions. If there is any problem, the transaction is rolled back.

Optimistic protocol: phases

- Read: reads all the elements in read-set and computes the result in local memory.
- Validate: validates the transaction by comparing read and write sets with those of other transactions. If the validation fails, the transaction is rolled back.
- Write: transaction writes the values in the write set.

Sets maintained by Scheduler for validation

- **START**: set of transactions that have started, but not yet completed validation.
- **VAL**: set of transactions validated and not finished the writing.
- **FIN**: set of transactions that have completed in Phase 3.
- For each transaction, the scheduler maintains $START(T)$, $VAL(T)$ and $FIN(T)$ timestamp.
- **FIN** is frequently purged.

Validation Rules for T

When we are trying to validate T

- Suppose there is U
 - If U is in VAL or FIN; that is, U has validated.
 - $\text{FIN}(U) > \text{START}(T)$; U did not finish before T started.
 - $\text{RS}(T) \cap \text{WS}(U)$ not empty. Risky, So rollback T.
- Suppose there is U
 - U is in VAL, U is successfully validated
 - $\text{FIN}(U) > \text{VAL}(T)$, U did not finish before T entered its validation phase.
 - $\text{WS}(T) \cap \text{WS}(U)$ not empty. Risky, So rollback T.

Comparison of Three CC protocols

- Locks:
 - Space is proportional to the number of database elements locked.
 - Delays transactions but avoids rollbacks
- Timestamps:
 - Space is needed for timestamps (naïve implementation)
 - Similar to the lock table, we can store the timestamps.
 - Rollback problem (detects earlier than validation)
- Validation
 - Space is used for timestamps and read/write sets.
 - Roll back problem

Summary

Have studied C.C. mechanisms used in practice

- 2 PL
- Multiple granularity
- Tree (index) protocols
- Timestamping
- Multiversion
- Validation

System Structure

