


Neighbourhood processing



$$I_{\text{out}}[m_0, n_0] = I_{\text{in}}[m, n]; [m, n] \in N[m_0, n_0]$$

Useful 2D functions

Unit impulse function in 2D:

$$\delta[m,n] = 1 ; [m,n] = 0 \\ = 0 \text{ otherwise}$$

Unit pulse (Rect) function in 2D:


$$\Pi[m,n] = 1 ; |m| \text{ or } |n| \leq 1 \\ = 0 \text{ otherwise}$$

Neighbourhood (spatial) operations

- Output pixel value at p in the output image is determined by pixels in N_p in the input image
- This is a **filtering** operation
- Filter kernel is represented as a **mask** or **window**
- The operation performed can be **linear** or **nonlinear**

Filtering / Masking operation

Initial mask position



Direction of movement of the mask is **top left → bottom right** of the image (raster scan)

Image

Final mask position

Non-linear operations

Notations: input image (x), output image (y)

- **Ranking**

- $y[m,n] = \max .. \min ... \text{median} \{x[m,n]; [m,n] \in N[m,n]\}$

Neighbourhood of (m,n)

- **Extrema** – special case in ranking

- $y[m,n] = \min \{x[m,n]; [m,n] \in N[m,n]\}$

- or $= \max \{x[m,n]; [m,n] \in N[m,n]\}$

help find the local max or min

Linear mask operations

y : output image x : input image h : filter mask

$y(m,n) = h(m,n) * x(m,n)$; h convolved with the input x

$$y[m,n] = \sum_i \sum_j x[i,j] h[m-i, n-j]$$

- output pixel is a weighted combination of the neighbouring pixels.
 - Weight is determined by the filter kernel h i.e. mask coefficients
- mask size is always finite ➔ FIR filtering

Linear filtering

What is linear filtering?

Filtered output (pixel) is a linear combination of input pixels

Common examples:

- Averaging (straight, weighted) or low pass filtering
- Finite difference or high pass filtering
- High boost filtering
- Bandpass filtering

Where is linear filtering used?

Common applications:

- Enhancement
 - Deblur, sharpen
- Denoising
 - By averaging
- Feature extraction
 - Detect edges

Spatial averaging

- Output pixel value at $[i,j]$ = average of all pixels in $N[i,j]$
- Mask weights = $h[i,j]$
- Mask is chosen to be odd sized 3x3, 7x7,...etc.
 - Why?
- Visual effect: *smoothing*
 - Degree of smoothing \propto mask **size**
 - Type of smoothing determined by mask **shape**, and **weights**. $h(i,j)$

$h[-1,-1]$	$h[-1,0]$	$h[-1,1]$
$h[0,-1]$	$h[0,0]$	$h[0,1]$
$h[1,-1]$	$h[1,0]$	$h[1,1]$

$m \times m$ mask
representing $N(i,j)$

Straight averaging

- Equal mask weights (Ideal low pass filter)
 - For $m \times m$ sized kernel, $h[i,j] = 1/m^2 \quad \forall i,j$
 - Equivalent to a running average filter in 1D
- Equivalent to convolution with $\Pi[m,n]$ *2D rect function or a box function*
 - Can lead to problems along sharp edges

12	10	17
17	12	17
35	16	184

$$x(m,n) = 12$$
$$\Rightarrow y(m,n) = 35$$

Smoothing using averaging



Before



After



After repeated
averaging

Check: Is repeated averaging by a filter the same as filtering with a larger kernel?

Weighted averaging

Gaussian smoothing : Mask weights are samples of a Gaussian function (distant neighbours have less weights)

$$G(x, y) = \frac{1}{2\pi\sigma^2} \exp \{ -(x^2 + y^2) / 2\sigma^2 \}$$

- Used to obtain controlled smoothing
 - Small σ means less smoothing and retaining fine detail
- Mask size increases rapidly with σ

N_x

$x(m,n) = 12$

12	10	17
17	12	17
35	16	184

h $\sigma = 0.391$ pixels

1	4	1
4	12	4
1	4	1

$y(m,n) = ?$

Controlled smoothing – 2 other alternatives

There are other ways to have controlled smoothing



1. Apply averaging selectively

Example


- include a neighbouring pixel in the averaging only if it is in a specified range
- output pixel is valid only if it is in a specified range,
- etc.

Controlled smoothing...contd.

2. Use a rotating mask and detect homogeneous region around a pixel of interest and average this region only



.....



9 rotated versions of a 3x3 mask

x : pixel of interest

Non-local Mean filter

- Non-local Means [Buades et al. 2005] – proposed for denoising

$$g_{NL}[r] = \sum_{q \in I} w[r, q] I[q] \quad r, q \text{ are pixel locations;}$$

$$0 \leq w \leq 1; \sum w = 1 \quad I[r] \text{ is the pixel value at } r$$

- w is the weight based on similarity between pixel values in the neighbourhoods of r and q .

Basic strategy: similar neighbourhoods should lead to large weights and vice versa

NLM filter contd.


Popular choice for w is a Gaussian weighted (Euclidean) distance between r, q

$$w(r, q) = \frac{1}{K(r)} e^{-\frac{N(\|I(\Omega_r) - I(\Omega_q)\|^2, a)}{h^2}}$$

$$K(r) = \sum_q e^{-\frac{N(\|I(\Omega_r) - I(\Omega_q)\|^2, a)}{h^2}}$$

$N(\mu, \sigma)$ denotes Normal distribution

Ω_x denotes the neighbourhood about x



From Buades et al. 2005

Sample result from Buades *et al*

Noisy input



Denoised results



Gaussian
averaging



NLM

Other types of spatial domain filtering

- Other common linear filtering operations
 - High-pass
 - Band-pass

What are the mask coefficients $h[m,n]$?

Clue: Any spectrum can be divided into a lowpass band + highpass band

➔ An **allpass filter** can be thought of as a sum of LPF and HPF

Check: What is the *filter kernel* h_{AP} for an **allpass filter**?

HPF Filter derivation

- Use LPF as the prototype

HPF: $h_{HP}[m,n] = \delta[m,n] - h_{LP}[m,n]$


$$\begin{aligned} y_{HP}[m,n] &= x[m,n] * h_{HP}[m,n] \\ &= x[m,n] * [\delta[m,n] - h_{LP}[m,n]] \end{aligned}$$

Using distributive property of convolution, rewrite as


$$y_{HP}[m,n] = x[m,n] - y_{LP}[m,n]$$

where $y_{LP}[m,n] = x[m,n] * h_{LP}[m,n]$

Examples



After HPF →



Quilt

BPF filter derivation

BPF: $h_{BP}[m,n] = h_{LP2}[m,n] - h_{LP1}[m,n]$

$$y_{BP}[m,n] = x[m,n] * (h_{LP2}[m,n] - h_{LP1}[m,n])$$

$$y_{BP}[m,n] = y_{LP2}[m,n] - y_{LP1}[m,n]$$

Filter masks

$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$

LPF mask

$-1/9$	$-1/9$	$-1/9$
$-1/9$	$8/9$	$-1/9$
$-1/9$	$-1/9$	$-1/9$

HPF mask

What is the sum of mask coefficients?

What is their significance?

Filter applications

- LPF
 - (impulse) noise smoothing,
- HPF
 - extract edges/details
- BPF
 - edge enhancement, sharpen noisy images

Other filters

- LPF smoothes and HPF retains only details in images

What if we want to only sharpen an image?

- **Unsharp masking or high boost filtering**
 - Derived from printing industry practice
 - Performs high frequency emphasis

$$y[m,n] = \alpha x[m,n] - x_{LP}[m,n] = (\alpha - 1) x + x_{HP} ; \alpha \geq 1$$

$\alpha = 1$ results in HPF (y has only details)

$\alpha > 1$ indirectly controls the degree of sharpening

Filter maskscontd

$w = ?$

-1/9	-1/9	-1/9
-1/9	w	-1/9
-1/9	-1/9	-1/9

High Boost Filter mask

Examples – HPF vs HBF



Input : Cheetah



3x3



5x5



7x7



15x15

HPF results for
different mask sizes



Input : Cheetah

High Boost filter results:
Effect of α and mask size



3x3 mask with $\alpha = 1$ (HPF)




3x3 mask with $\alpha = 1.5$




3x3 mask with $\alpha = 2$



7x7 mask with $\alpha = 1.3$



Santa Fe -
Remote sensing



Santa Fe - HPF



Santa Fe - HBF

Examples



Face



Face - HPF



Face - HBF

Non-linear filtering – some applications

Median filtering


- Output pixel = median value of all pixels within the $N(i,j)$
- *Effects*
 - Excellent in removing salt and pepper noise
 - Does not shift boundaries
 - Does not affect brightness change across steps
- Computationally intensive as processing an image with M pixels requires M sorting operations

Median filtering - example

12	10	17
17	12	17
35	16	184

Pixels within a 3x3
window

Ranked list: 184 35 17 17 17 16 12 12 10



output pixel value = Median = 17

Comparison of different filters

Input image patch

12	10	17
17	12	17
35	16	184

Output pixel value

Mean filter: 35 **Gaussian weighted mean filter : 28**
Median filter: 17

Mean after ignoring the 'noisy pixel' = 17

Application 1: Reducing noise



Noisy original



After averaging



After median
filtering

Application 2: Finding extrema

Task: find extrema points in an image.


Solution 1 **Max (Min) filter**

- Output pixel at $p = \max\{I_{\text{in}}[m,n]; [m,n] \in N_p\}$
 - Simple ranking
 - Does not discriminate between extrema of different kinds

Solution 2 **Top (bottom) hat filter**


- Output pixel at $p = I_p$ if $I_p >$ pixel level in its background
 - *Context-aware* max finder
- Uses 2 neighbourhoods (1 each for estimating background and foreground)

Top hat filter



1-D case: finds local peaks/maxima

- Low or broad peaks are skipped



Top hat function

For the local peak of interest
the control parameters are:

w_h : peak *width*

w_h : specifies the *extent* of
the context

h : peak *height*

Top hat filter


- “Interesting” point finder
 - w_h specifies the spread of the peak of interest
- Many versions exist

Example: to find the bright points in an image


A method:

1. Define a central and a surround region about a pixel (with a window)
2. Rank the pixels in the two regions and find the max values
3. Find $g_{cmax} - g_{smax} = g_{diff}$
4. If g_{diff} is above a threshold **retain** centre pixel.
Else set it to 0

Neighbourhood patterns



Octagonal



Square

Top hat filtering

Effects:

- Only “interesting pixels” are retained
 - is a suppression operator
- Centre pixel is retained even if it is not the brightest pixel (as the criterion is $g_{\text{diff}} > t$)
 - produces small “halo” around bright points
- Matlab also has a *Bottom hat* function


Stars in Night Sky



Top Hat Filter Result



Max Filter Result




Lena Image



Top Hat Filter Result



Max Filter Result



Neighbourhood processing – implementation issues

Issues

Handling border pixels: what are their neighbours?

Option 1

- Don't compute the output for such pixels
 - Output image is smaller than input image

Option 2

- Extend the image size by zero padding before processing

Option 3

- Extend the image by one of many ways (replicate, wrapped around, etc.) before processing

Implementation issues.. contd.

Computational load – load increases with mask size as well as type of filtering

- Pixel addressing load as well as processing load

Solution 1

- Implement in frequency domain
 - Exploit the property that convolution in spatial domain is equivalent to multiplication in frequency domain

Check: what is the max kernel size above which it is not efficient to implement filtering in spatial domain?

Solution 2

- efficient hardware implementation

When to use what type of proc.?

Point

- When type of image is *fixed*
 - Intensity characteristics of object/region of interest is *known*
- Ex. Medical images of specific organs
- One time parameter tuning

Neighbourhood

- More *generic* and *customizable*
- *Demands skill* in modeling the neighborhood interaction

Global

- When type of image is *unknown*
- Objects/regions of interest are *unknown*
so intensity characteristics
- i.e. Natural images/indoor/outdoor