

Image transforms

Why transform?

- simplify some operations
- compactly represent given image
- enable easy extraction of information from an image

Background

- Inner product : $\langle x, y \rangle = \sum_j x_j y_j^*$
- Any 2D vector v can be expressed as a linear combination of a basis $B: \{e_j\}$

$$v = \sum_j \alpha_j e_j;$$

$$\alpha_j = \langle v, e_j \rangle$$

- If e_j are unit length and mutually orthogonal, B is an orthonormal basis (ONB)

Background

Identity matrix = $\{\delta[m-n]\}$

$A^H = A^{*T}$ (conjugate transpose)

Real case

A is symmetric : $A = A^T$;

A is orthogonal: $AA^T = I$ or $A^{-1} = A^T$

Complex case

A is Hermitian : $A = A^H$

A is unitary: $AA^H = A^H A = I$ or $A^{-1} = A^H$

A is real and orthogonal \rightarrow A is unitary (Converse is **not** true)

Image transforms – decomposition view

- Image transforms are a way to decompose images in terms of **basis** images
- A given image can be expressed as a weighted combination of basis images
- The weights are found as the inner product of the given image x and the basis image ϕ .

Image decomposition $x[m, n] = \sum_{l=0}^{N-1} \sum_{k=0}^{M-1} X[k, l] \phi[m, n, k, l]$ Synthesis

Weights
or transformed image $X[k, l] = \langle x[m, n], \phi[m, n, k, l] \rangle = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} x[m, n] \phi^*[m, n, k, l]$ Analysis

Fixed basis: Fourier, DCT, Slant Haar, Hadamard, Transforms

Non-fixed basis: Wavelet

Image transforms – matrix view

Any image x can be transformed to X by pre- and post-multiplying x with matrices

$$\begin{aligned} X &= PxQ && \text{analysis} \\ \therefore x &= P^{-1}XQ^{-1} && \text{synthesis} \end{aligned}$$

We can also write these using double summations:

$$\begin{aligned} X[k, l] &= \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} P[k, m] x[m, n] Q[n, l] \\ x[m, n] &= \sum_{k=0}^{M-1} \sum_{l=0}^{N-1} P^{-1}[m, k] X[k, l] Q^{-1}[l, n] \end{aligned}$$

Choice of P and Q matrices

P and Q are need to be chosen such that

- inversion of the transformation is possible

$$\triangleright x \Leftrightarrow X$$

- transform (X) computation is efficient
- X reveals some information about x

Types of transforms

- If $\phi(m,n,k,l) = \phi_1(m,k) \phi_2(n,l)$ then we have a **separable** transform

- If $\phi(m,n,k,l)$ is a set of orthogonal basis then we have a **orthogonal** transform

$$\langle \phi[m,n,k,l], \phi^*[m,n,k',l'] \rangle = \delta(\overset{\downarrow}{k} - \overset{\downarrow}{k'}, \overset{\downarrow}{l} - \overset{\downarrow}{l'})$$

- If $\phi(m,n,k,l)$ is a complete set then

$$\sum_{k=0}^{M-1} \sum_{l=0}^{N-1} \phi[m,n,k,l] \phi^*[m',n',\overset{\downarrow}{k},\overset{\downarrow}{l}] = \delta(\overset{\downarrow}{m} - \overset{\downarrow}{m'}, \overset{\downarrow}{n} - \overset{\downarrow}{n'})$$

- If all the above hold, $P=Q$ and we have a **separable, unitary** transform

$$X(k,l) = PxP^T \text{ and } x(m,n) = ?$$

Discrete Fourier Transform

DFT: $x \rightarrow X$; with a ONB made of complex sinusoids

- Elements of P and Q matrices are **complex** exponentials

If x is of size $M \times N$, then P is $M \times M$ and Q is $N \times N$

$$P[k, m] = \frac{1}{M} e^{-j \frac{2\pi}{M} km} \quad Q[n, l] = \frac{1}{N} e^{-j \frac{2\pi}{N} nl}$$

- X is a complex valued function
 - Two images are needed to display X : {amplitude and phase}
or {real and imaginary}

Note: $P = Q$

Discrete Fourier Transform

Let $P = Q =$ $F(k, m) = e^{-j\frac{2\pi}{N}mk} = W_N^{-mk}$


Then $X = F x F$ and $x = F^{-1} X F^{-1}$

- F is unitary and symmetric
➤ $F^{-1} = F^H$ and $F = F^T \rightarrow F^{-1} = F^*$

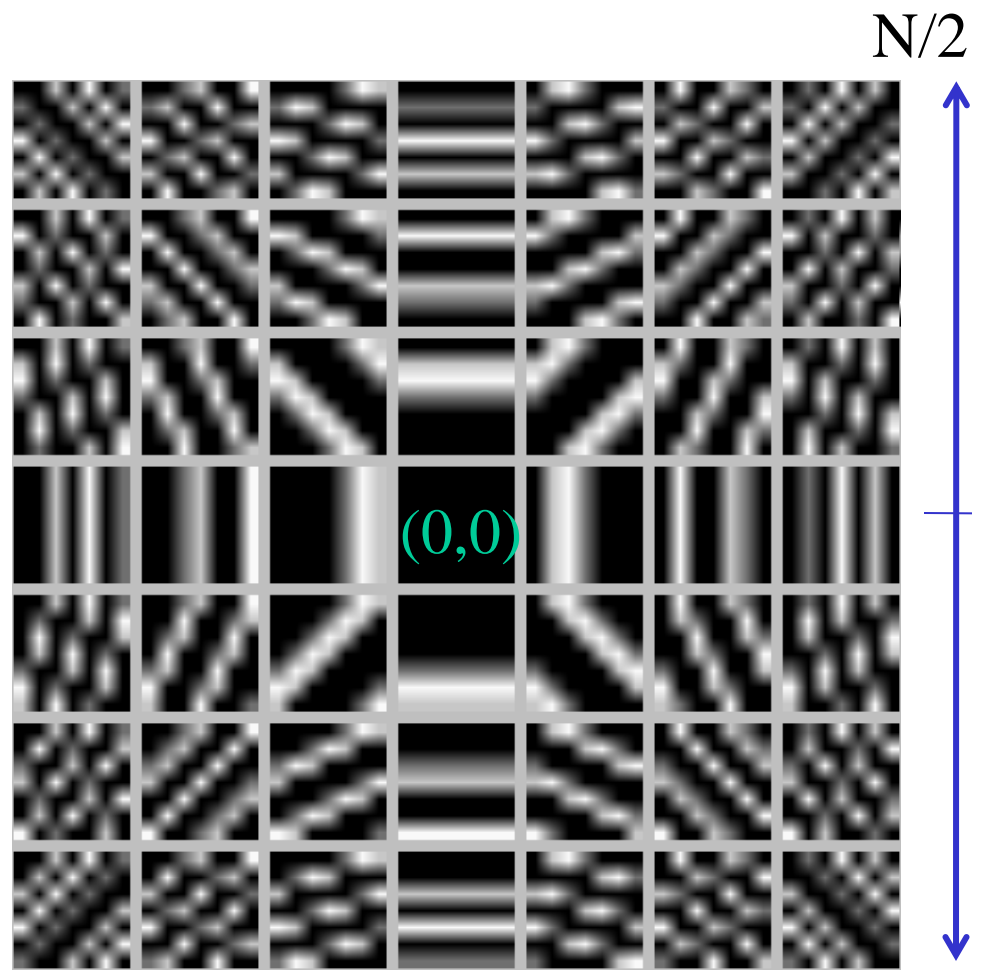
- Basis images for a $N \times N$ image :

$$\left\{ \frac{1}{N} W_N^{(mk+nl)} \right\} \quad 0 \leq k, l, m, n \leq N-1$$

Basis Images of DFT



Real part



Imaginary part

Symmetry properties of DFT

- The basis function is periodic (with N)

$$X(k) = X(k \pm N)$$

$$X(k + aN, l + bN) = X(k, l)$$

$$0 \leq k, l, \leq N-1$$

- Conjugate symmetry about N/2

$$X\left(\frac{N}{2} - k\right) = X^* \left(\frac{N}{2} + k\right)$$

$$X\left(\frac{N}{2} \pm k, \frac{N}{2} \pm l\right) = X^* \left(\frac{N}{2} \mp k, \frac{N}{2} \mp l\right)$$

$$X(k, l) = X^*(N - k, N - l)$$

$$0 \leq k, l, \leq \frac{N}{2} - 1$$

➤ There are only N Unique values for X


DFT computation

- When one computes a $N \times N$ DFT for an image x , the result does not capture 1 complete period of the spectrum if origin is set at top left corner
- Hence, to display **one entire period of the spectrum** we **shift the origin to the centre**
- To do this we note $(-1)^{m+n} x(m, n) \Leftrightarrow X(k - \frac{N}{2}, l - \frac{N}{2})$



Checker board pattern

Computing and displaying $X[k,1]$



Computed DFT



Does not help visualise 1-full period of the spectrum

Displayed DFT

Covers 1-full period of the spectrum

DFT interpretation

DFT with origin shift Computed $N \times N$ DFT



Violin image




DFT (amplitude) spectrum

Violin image is periodically extended!



It is periodic!

DFT spectra of images



$|X|$: Magnitude spectrum
 $\angle X$: Phase spectrum

Transform

Amplitude spectrum

The amplitude spectrum of an image typically has


- A very high dynamic range
 - Well above the 256 levels of a 8-bit display
- Energy in LF band > energy in MF band > energy in HF band
- A non linear greyscale transformation is needed to display properly
 - Ex. $\text{Log} (1 + |F|)$; F is the Fourier transform of image f

Check: why log transform?


Importance of non-linear stretching



Piano



$|\text{Piano}|$



$\log(|\text{Piano}|+0.01)$

DFT examples ...

Raw Image

Fourier Amplitude

Horizontal
Line





Vertical Line

Vertical Line



Horizontal
Line

Diagonal Line





Diagonal Line

DFT examples

Raw Image


Fourier Amplitude

Vertical sinusoid,
higher frequency



DC term +
side lobes
wide spacing

Vertical sinusoid,
lower frequency





DC term+
side lobes
close spacing

DFT examples ...

Raw Image


Fourier Amplitude

Horizontal sinusoid,
higher frequency



DC term +
side lobes
wide spacing


Sinusoid,
tilted




Tilted spectrum

DFT examples: (symmetric & binary) ellipse


ellipse





|ellipse|





\angle ellipse



Rotated ellipse





Amplitude spectrum examples: (binary) rectangle




A rectangle and its $|FT|$

Amplitude spectrum examples: (binary) circle



A circle and its |FT|



DFT of a periodic image





|chess|

origin is at top left


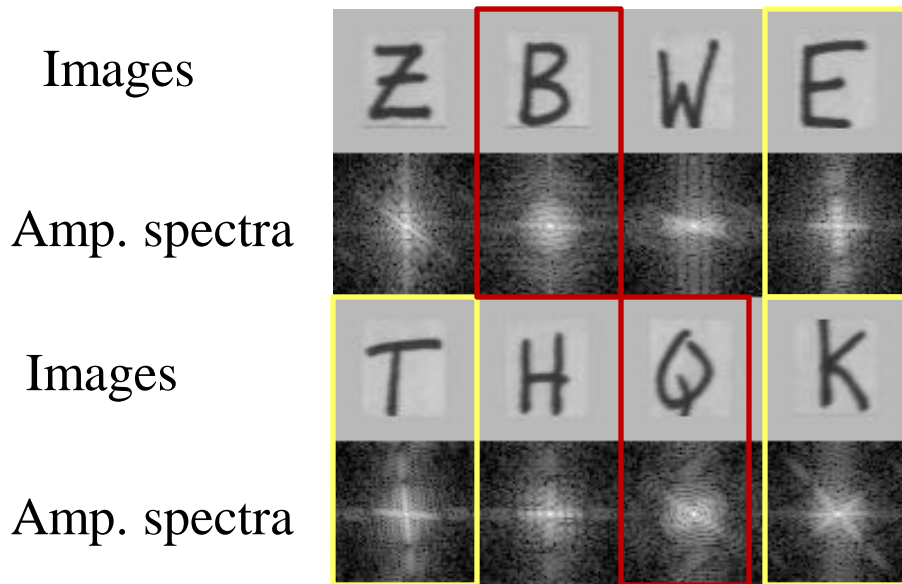
∠ chess




with origin shifted to the centre




DFT of greyscale images - effect of strong edges




DFT examples: (greyscale image) Ape



Ape




|ape|




\angle ape

Inverse DFT with incorrect phase spectra

IDFT $\{|\text{ape}|, \angle \text{ellipse}\}$



IDFT $\{|\text{ellipse}|, \angle \text{ape}\}$



The phase spectrum appears to be strongly encoding the shape

Filtering in the Transform domain

Linear Filtering

- Given $x[m,n]$ the filtered image $y[m,n]$ is

$$y[m,n] = (x * h)[m,n]$$



$$y[m,n] = \mathfrak{F}^{-1}\{Y\} = \mathfrak{F}^{-1}\{X[k,l]H[k,l]\}$$

where $H[k,l] = DFT\{h[m,n]\}$

- H can be viewed as a frequency domain **mask**
- Mask is *binary valued* for *ideal* filtering and *grey valued* for *non-ideal* filters (ex. Gaussian)
- Size of H and X are equal
 - Y is the result of multiplying X and H

Convolution \Leftrightarrow product

An important property of Fourier transform:

$$\mathfrak{F}\{f[m,n] * h[m,n]\} = F[k,l]H[k,l] \quad (1)$$

- * here is a periodic convolution since DFT and IDFT are periodic
- So (1) will result in wraparound errors if we want linear convolution

Implementing linear convolution using DFT:


Let f be $M \times M$ and h be $L \times L$

\therefore linear convolution of f and h will result in an image of size $M+L-1$


Extend f and h to be $(M+L-1) \times (M+L-1)$ by zero-padding

before implementing (1)

Test image1 and its DFT




rings




|rings|

Test image2 and its DFT



star




|star|

Low pass filtering

- Ideal/brickwall filter**

$$H(u,v) = 1 ; \omega = \sqrt{u^2 + v^2} < \omega_c = \sqrt{u_0^2 + v_0^2}$$
$$= 0 \text{ otherwise}$$


ω_c is the cutoff frequency



- Non ideal filter - Butterworth filter**


$$H(u,v) = \frac{1}{\sqrt{1 + \left(\frac{\omega}{\omega_c}\right)^{2n}}}$$

n is the order of filter




u-v plane binary mask

Low pass filtered rings




Ideal filter




Butterworth filter

Low pass filtered star




Ideal filter



Butterworth filter


High pass filtering

(Ideal) $H(u,v) = 1 ; (u^2+v^2)^2 > (u_0^2+v_0^2)^2$
 $= 0$ otherwise




u-v plane binary mask

High pass filtered ring




Ideal HPF




Butterworth filter

High pass filtered star




Ideal filter




Butterworth filter

Filtering with zonal masks

- Image is a 2D signal → we have additional degree of freedom in filtering
- Can design filters to select image features at specific **spatial frequency and orientation**
- Implementable using masks of different shapes



Fan filter



Wedge filter

Some applications

Correlation

Correlation operation

$$f[m, n] \circ g[m, n] = y[m, n] = \frac{1}{M^2} \sum_{i=0}^{M-1} \sum_{j=0}^{M-1} f^*[i, j] g[m+i, n+j]$$

conjugate
↙

Popular as a template matching technique

- To find similarity between an unknown and a set of known images
- To locate instances of a given image within a larger image (ex. find an alphabet in a document image)

In the Fourier (Frequency) domain

$$\mathfrak{F}\{f[m, n] \circ g[m, n]\} = \bar{F}[k, l] G[k, l]$$

Correlation can be computed with DFT after zero-padding f and g

Illumination correction

Image model: $f[m,n] = i[m,n] r[m,n]$

i – illumination (corrupted) r – reflectance (of interest)

Filtering steps

1. Convert the multiplication model to additive using a *log* transform

In the transformed space

2. Estimate the degradation $i[m,n]$
3. Subtract from given $f[m,n]$
4. Find the inverse of the *log* transform to recover the corrected image

Homomorphic filtering

Filtering steps

1. Do a Log transform of f : $\ln f[m,n] = \ln i[m,n] + \ln r[m,n]$
I.e. $x[m,n] = a[m,n] + b[m,n]$
2. Compute the Fourier Transform $X[u,v] = A[u,v] + B[u,v]$
low pass high pass
3. **Lowpass filter** X to extract the illumination component
 $LPF \{ X \} \sim A \rightarrow IDFT\{A\} = a$
4. Perform illumination correction: $y = x - a$
5. Perform inverse Log transform of y to obtain the desired illumination corrected image: $e^{(x-a)}$


Homomorphic filtering - effects

- Corrects illumination but reduces the dynamic range
 - Lowers the brightness level of entire image
- Variable results as f is only estimated




Homomorphic filtering – another example

Original



Filtered



Decomposition of an image

So far we looked at

Discrete Fourier Transform

- decomposes images using a fixed basis
 - complex exponential
- What if we use a non-fixed basis?

Examples:

1. The **SVD** decomposition

- which uses an **image-adaptive** basis
 - the basis is derived from the given image

2. The **Wavelet Transform**

- which uses a non-image adaptive, non-fixed basis

SVD - Diagonalising a given image

Claim: Given any image $f(m,n)$ we can diagonalise it (make it sparse)

Let $g = ff^T$.

The eigenvectors \mathbf{u}_i of g are found from: $g\mathbf{u}_i = \lambda_i\mathbf{u}_i$

The eigenvectors \mathbf{v}_i of g^T are found from: $g^T\mathbf{v}_i = \lambda_i\mathbf{v}_i$

Let $U = \{\mathbf{u}_i\}$ and $V = \{\mathbf{v}_i\}$


Then, the **Singular Value Decomposition (SVD)** of f (of rank r) is

$$f = \sum_{i=1}^r \lambda_i^{1/2} \mathbf{u}_i \mathbf{v}_i^T \quad \text{Or in matrix form} \quad f = U\Lambda^{1/2}V^T$$


Λ is a diagonal matrix of r non-zero eigenvalues of g

SVD based image recovery

Original image 256 singular values




retaining 20 singular values




Good solution for
compression ?

retaining 50 singular values



retaining 85 singular values



SVD of images

- If Λ is constructed with decreasing λ_i then SVD is an **optimal decomposition** for an image f
- f can be approximated as a linear sum of a small set (b) of basis images with least square error
- Basis images of SVD are eigenimages of f found as $\mathbf{u}_i \mathbf{v}_j^T$

Note: SVD is similar/closely related to Karhunen
Loeve Transform (KLT) and Principal
Component Analysis (PCA)

To learn more, see Gerbrands JJ (1981) *On the relationships between SVD, KLT and PCA*, **Pattern recognition**, Vol. 14, No. 1-6. pp. 375-381.