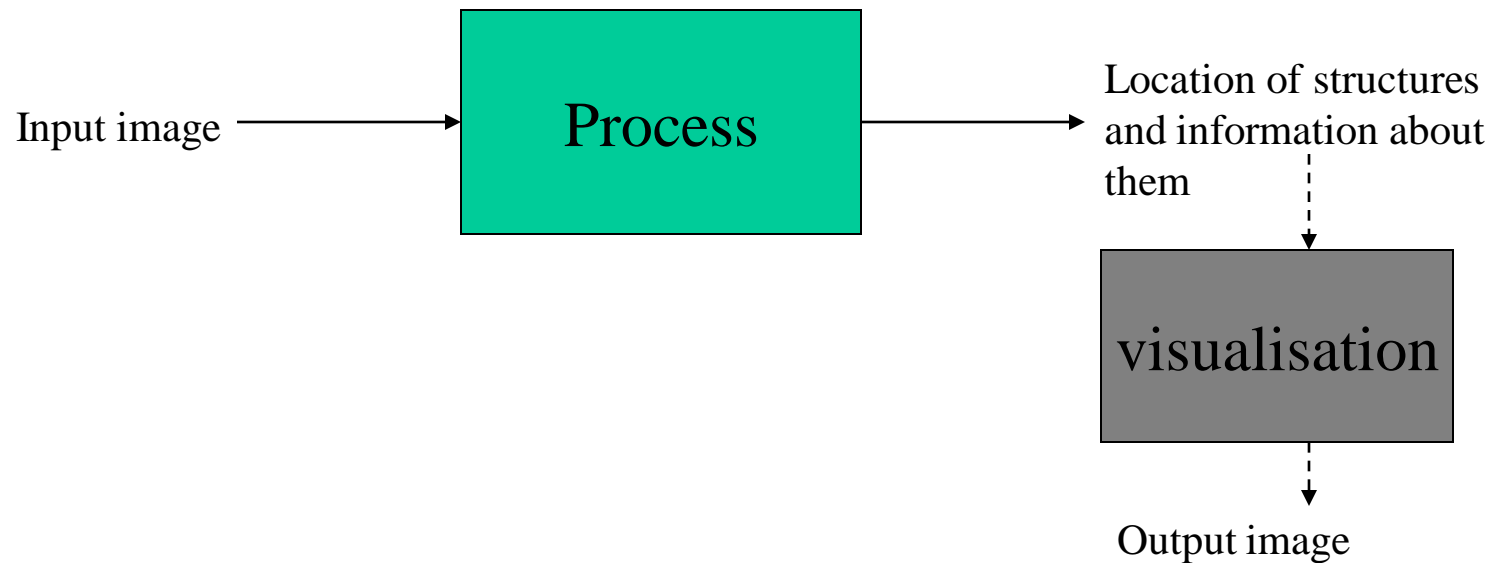


Image Analysis

What is image analysis?

- Processing images to organise the image into structures and derive information about them



Example application areas

- Image measurement
 - Area, perimeter, etc.
 - Ex. remote sensing
 - No of objects
 - Ex. medical images
- Object recognition
 - Faces, animals, buildings, license plates, etc
 - Terrain types
 - Lesions, tumors


Key problems in image analysis

- Segmentation
- Feature detection
 - Feature extraction
 - Shape analysis
 - Texture analysis
- Motion analysis

Segmentation

Goal: Organise the image into meaningful groups/regions

- Groups are characterised by content




16 distinct regions/groups

How many regions/contours in this image?



Tough!

Object recognition



We see '*objects*' by grouping basic elements (blobs/lines)

Rule: Elements are **neighbours** *and* they are **semantically related**

- different ways of grouping → different objects!

Grouping in Human Vision

How do we know which parts of visual input belong together ?




- This has been studied by perception theorists
- A set of principles have been identified:

Max Wertheimer's concept of *pragnanz*:

“when things are grasped as wholes, minimal amount of energy is exerted in thinking”

Key principles

- **Emergence** – simple grouping rules lead to complex pattern formation
- **Reification** – perception is constructive/generative
- **Multistability** – multiple percepts can be stable and switch back and forth
- **Invariance** – recognition of simple objects is invariant to geometrical transformations







Gestalt principles of grouping

- **Proximity** – group based on neighbourliness of blobs
- **Common fate** - group if there is coherent motion
- **Parallelism**- group parallel curves/lines
- **Closure** – group if it leads to closed figures
- **Continuity** – group if continuous in space or feature
- **Similarity**- group based on some shared feature
- **Symmetry** –group if it can result in symmetric structures
- **Familiar Configuration** – if blobs when grouped, lead to a familiar object do the grouping!

cognitive

Grouping examples



Why is segmentation hard?

- It is an ill-posed problem!
- For natural images, the benchmark (ground truth) is human perception
- Need to make the problem tractable

Segmentation – towards a formulation

Goal of segmentation is to organise the image into groups/regions such that there is

1. Homogeneity within regions

- Interiors are devoid of holes

2. Distinctness across adjacent regions

Segmentation – formulation 1

Given image $x[m,n]$,

$$x[m,n] = \bigcup_i R_i \quad ; \bigcap_i R_i = \phi$$

Find all the pixels belonging to a segment:

Label each pixel as belonging to some R_i

- Membership (unique) in a region is based on **shared property** (homogeneity within a region)

Segmentation - formulation 2


Find all points that mark the interface between segments:

- $x[m,n] = \{c_i\}$ a set of contours or boundaries
- A contour point is a point where x is *discontinuous* in some feature space
 - Distinctness across regions

Segmentation into 2 classes

2 class case (binary output)

2 - class pixel classification



Thresholding - issues

- Division of image into 2 **uniform** regions

- Object vs background

- Uniformity in intensity is popular

if $x[m,n] > t$

$x[m,n] = 255$ *object*

else

$x[m,n] = 0$ *background*

Problems:

- Thresholding (aka binarisation) is a point processing method
 - Context and spatial relations are ignored
 - Contiguous regions are not guaranteed
 - Will be sensitive to noise
- Post proc. is required for corrections

Thresholding approaches

Automatic thresholding

1. Fixed

- limited use as it is sensitive to any change in image characteristic
- seen earlier

2. Adaptive

- Global
- Local
- Dynamic

Adaptive thresholding - Simple Methods

Adaptive thresholding

x : given image

p_l : some local property

t : threshold

- **Global:** $t = f\{x[m,n]\}$
global property

ex: $t = \alpha \mu_x$
global mean

- **Regional:** $t = f\{p_l\}$
local property

ex: $t = a \mu_l$

- **Dynamic:** $t = f\{[m,n], p_l\}$
location local property


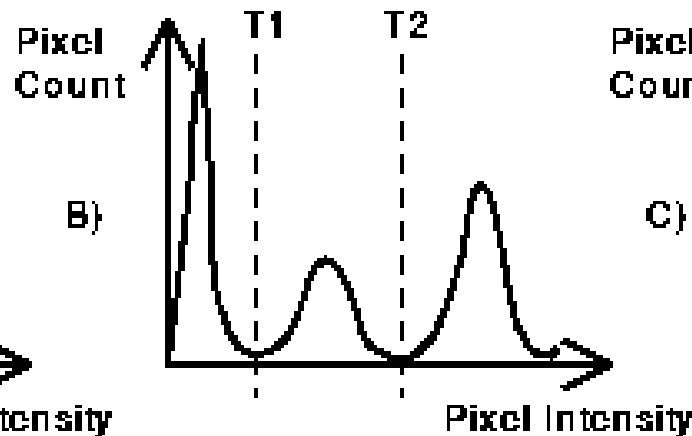

ex: $t = a \mu_x + b \sigma_{mn}$
local standard deviation

Global thresholding

- Applied to the whole image
- Uses image data to determine the threshold
 - Brightness histograms $h[g]$ are popular
 - h should have well defined peaks and troughs for best results
 - Single or dual thresholds {trough} or {trough 1: trough 2}
- Computationally less intensive

Selecting the threshold

Number of peaks and definition of troughs influence threshold selection



Difficult case

Histogram smoothing

- When a histogram is ragged, it can be smoothed before selecting thresholds
- Threshold selection is made easier

Warning: Smoothing can shift peaks

➤ Depends on size of smoothing kernel

Regional thresholding

- Divide the image into patches prior to thresholding
- Patch boundaries need to be post processed
- Combinations of local statistics and other derived information are used to find t
 - ex. Texture, entropy etc.
- Effective for handling non-uniform illumination
- Computationally more intensive

Thresholding- Example

Scanned doc



b to f. Results of binarisation with different methods

Thresholding colour images $t = 127$



Where is the face?

Thresholding colour images $t = 127$



on green

on blue



Where is the bird?

Iterative and non-iterative methods

- Many traditional methods have been proposed
 - Triangle
 - Isodata
 - Otsu
- No universal solution exists!
- Newer more complex methods continue being proposed

Triangle algorithm


- A non-iterative method

Algorithm

From a given image with histogram $h[g]$

1. Find the line l between the highest peak
i.e. $\max\{h[g]\}$ and g_{min}
2. For every g find distance $d(g)$ from l to $h[g]$
3. Desired $t = \operatorname{argmax} d(g)$

Triangle algorithm



Effective when objects produce weak peaks

Isodata algorithm

- An iterative method

Algorithm:

1. Start with threshold $T_0 = 2^b - 1$ (for a b -bit grey scale image) to get 2 regions, R_1 and R_2

For $k > 0$

1. $T_{k+1} = 0.5 (m_{k,1} + m_{k,2});$
 $m_{k,1}$ = mean grey value of R_1
2. Iterate until $T_{k+1} = T_k$

Note: No assumption is made about the distribution of pixel values in object/background

Otsu's method

$h_x(g)$: the histogram of image $x[m,n] \rightarrow p_x(g)$ pdf


$g \in [0, L-1]$

o : object b : background

Task: need to find a threshold $t = g_o$ from h_x

Let $g < t$ belong to b class and $g \geq t$ belong to o class

Need to minimise error probability in classification –
MAP detector



Otsu's solution:

optimal t is one which
minimises the variance
within each class

$$\sigma_{within}^2(t) + \sigma_{between}^2(t) = \sigma^2$$

Otsu's solution ..contd.

$$\sigma_{between}^2(t) = n_b(t)\sigma_b^2(t) + n_o(t)\sigma_o^2(t) \quad \text{Weighted sum of class variances}$$

$$n_b(t) = \sum_{g=0}^{t-1} p(g); \quad n_o(t) = \sum_{g=t}^{L-1} p(g)$$

Similarly $\mu = n_b(t)\mu_b(t) + n_o(t)\mu_o(t)$

We can show that

$$\sigma_{between}^2(t) = n_b(t)n_o(t)[\mu_b(t) - \mu_o(t)]^2 \quad \text{Weighted squared difference of class means}$$

Derivation for inter-class variance

$$\begin{aligned}
 \sigma_{between}^2(t) &= \overset{\text{global}}{\sigma}^2 - \sigma_{within}^2(t) \\
 &= \left[\frac{1}{N} \sum_{m,n} x^2[m,n] - \mu^2 \right] - [n_b(t)\sigma_b^2(t) + n_o(t)\sigma_o^2(t)] \\
 &= \left[\frac{1}{N} \sum_{m,n} x^2[m,n] - \mu^2 \right] - [n_b(t) \sum_{m,n} x^2[m,n] - \mu_b^2] - [n_o(t) \sum_{m,n} x^2[m,n] - \mu_o^2] \\
 &= n_b(t)[\mu_b(t) - \mu]^2 + n_o(t)[\mu_o(t) - \mu]^2; \quad \text{N: number of pixels in } x \\
 &= n_b(t)n_o(t)[\mu_b(t) - \mu_o(t)]^2 \text{ since } \mu = n_b(t)\mu_b(t) + n_o(t)\mu_o(t)
 \end{aligned}$$

Otsu's algorithm

1. For every t_k , threshold x and bin the results into 2 bins.
2. Compute $\sigma_{between}^2(t_k) = n_b(t_k)n_o(t_k)[\mu_b(t_k) - \mu_o(t_k)]^2$
 n_b and n_o are number of pixels in the 2 bins
3. Required $t = \arg \max \sigma_{between}^2(t_k)$

Can be implemented recursively

Comparison of thresholding methods

- There are many more methods to find threshold
- For a survey check

Sezgin, M and Sankur, B (2004), "Survey over Image Thresholding Techniques and Quantitative Performance Evaluation", Journal of Electronic Imaging 13(1): 146-165

- For a comparison check

http://www.fmwconcepts.com/imagemagick/threshold_comparison/index.php

Segmentation– region based approaches

Region-based approaches

- Grouping by similarity and spatial proximity
- Suitable for segmenting non-uniform regions that are perceived to be uniform
- Region growing algorithms



Pixel aggregation

Split & merge

Pixel aggregation

- A **bottom-up** approach to segmentation

Algorithm:

1. Start with a seed pixel
2. Add neighbouring pixels if they are “similar”
 - Similarity in grey value, texture, moments, colour etc
3. Stop when no new pixels are added

Issues in pixel aggregation

Seed selection

- Manual or automatic
- *A priori* knowledge is essential for good results

Criteria for aggregation

- Pixel value, local statistics, connectivity
- Too small vs large area for testing

Stopping the growth


- Change in size, shape, some knowledge about boundary
- ❖ Inefficient for detecting large number of regions in a complex image
- ❖ Results are sensitive to *seed choice* for inhomogeneous regions
- ❖ To find K segments need to do K region growing operations

Region growing example

Input image



methods



Human marked segments



Computed segments

Quad tree (1971)

Given an image f of size $(2^k \times 2^k)$, recursively divide it into smaller regions

Method:


Define a measure χ of intensity variation

Set $f = f^k$

1. If $\chi(f^k) > \alpha$ then split f^k into subimages $f_j^{k-1}; j = 1, \dots, n$
2. Repeat previous step on f_j^{k-1}

Final result: a tree of degree n , leaves are homogeneous subimages

Splitting using quad tree



Merge

Merging criterion: homogeneity

Method: Scan the split results and check for homogeneity of adjacent regions and merge them

Issues in split and merge

Splitting method

- Use of quadtrees - Can result in blocky boundaries

Criterion for splitting

- Test for inhomogeneity – local statistics
 - Standard statistical tests assume normal distributions which is rarely true

❖ Quality of result depends on testing criteria

❖ All segments can be found in one run

❖ No. of iterations depend on image content and criterion choice

Example

9	10	10	10	10	10	10	10
5	50	55	60	50	50	50	9
10	55	52	55	200	10	55	9
10	50	200	50	54	5	55	10
10	60	200	200	54	57	60	10
10	58	10	10	10	50	58	10
10	52	55	60	55	60	60	9
10	9	9	9	9	9	9	10

$g_{\text{mean}} = 55$; $\text{delta} = \pm 5$ Red colour – ideal segment

Segmentation: Probabilistic approaches

Relaxation

- An **iterative** pixel labeling approach
- Unlike histogram-based approach, it takes into account both greyvalue and context of a pixel
 - Can incorporate local constraints in labelling

Uses a probabilistic reasoning for assigning labels

- p_{xl} : probability that pixel x has a label l
 - p_{xl} also depends on labels of neighbouring pixels
 - To enforce homogeneity criterion

Probabilistic relaxation - method

pixel values $f_i ; i = 1, 2..N$

classes $l_m ; m = 1, 2..M$

For every pixel pair f_i, f_h with labels l_j, l_k

Define

1. a compatibility measure $C(f_i, l_j ; f_h, l_k)$
 $C > 0 \rightarrow$ labels are compatible and vice versa
 $C \sim 0 \rightarrow$ uncertain compatibility
2. $p_{ij} :$ probability that f_i has label l_j

Strategy:

- initialise the probabilities
- if p_{hk} is high and $C(f_i, l_j ; f_h, l_k) > 0$ then increment p_{ij}
- if p_{hk} is high and $C < 0$ then decrease p_{ij}
- if p_{hk} is low OR $C(f_i, l_j ; f_h, l_k) \sim 0$ then do nothing to p_{ij}

Compatibility function

Practical assumptions:

- Defined only for neighbouring pixels
- It is spatially invariant
 - Requires $8M^2$ computations for a 3×3 neighbourhood

Example – detect smooth curves in an image

- Use the slope (θ_j) at every pixel f_i to define the initial p_{ij} ;
- p_{im} is the probability there is no curve at f_i
- Definition for C

$$c(i, j; h, k) = | \cos(\theta_j - \theta_{ih}) | \cos(\theta_k - \theta_{ih}) |$$

$$c(i, m; h, k) = -\cos 2(\theta_k - \theta_{ih})$$


- $c(i, j; h, k) \rightarrow C = 1$ for collinear points and $C = 0$ for ?
- $c(i, m; h, k) \rightarrow$ a curve with slope θ_k at f_h is **incompatible** with no curve at f_i .

Probabilistic relaxation


- **Adv:** handles noisy conditions well
- **Disadv:** can suffer when structures have non-uniform shapes

Example


Input




Results



Thresholding



Relaxation



Relaxation - extensions

Relaxation labelling technique has evolved over time

- a pixel is represented by a host of features (not just intensity)



Newer techniques are based on :

- Expectation maximisation (EM)
- Gaussian mixture modelling (GMM)

Naïve Bayes theorem

x: pixel c: class

$$p(c | x) = \frac{p(x | c) p(c)}{p(x)}$$

$p(c/x)$ *a posteriori* probability

$p(x/c)$ likelihood

$p(x)$, $p(c)$ are *prior* probabilities or ‘Priors’

Segmentation as Expectation Maximisation

Assumption: Pixels in class l can have different values
→ represent it with a feature vector

θ_l : a vector of parameters (mean, variance of greyvalues, texture, etc) associated with pixels in class l

Define

- $p_l(x/\theta_l)$; $l = 1, 2..L$ the probability distribution of a pixel given a class label

EM formulation: Segmentation is an incomplete data estimation problem

- *incomplete* data are the measured feature vectors θ_l

EM algorithm

- Helps find the Maximum likelihood (ML) or Maximum a posteriori (MAP) estimate

EM algorithm

Initialize an estimate of the parameter vector $\theta_l, l = 1, 2..L$

Repeat

1. E step

Estimate the labels based on the current parameter estimates

2. M step

Update the parameter estimates based on the current labelling

Until Convergence

Gaussian Mixture Modelling and EM

- Assume the probability distribution of the pixels in different classes to be **Gaussians**
- For L classes we have a mixture of L Gaussian distributions

$$p(x | \Theta) = \sum_{l=1}^L \alpha_l p_l(x | \theta_l) \quad \text{Mixture of Gaussians}$$

$$\Theta = \{\mu_l, \Sigma_l, \alpha_l\}$$

$$p_l(x | \theta_l) = \frac{1}{\sqrt{2\pi} \det(\Sigma_l)^{1/2}} e^{-\frac{1}{2}(x-\mu_l)^T \Sigma_l^{-1} (x-\mu_l)}$$

E step now has a linear solution

GMM based segmentation

The E step :

Given a pixel x_j and its parameter vector θ_j
find its label l as

$$P(l | x_j, \theta_l) = \frac{\alpha_l p_l(x_j | \theta_l)}{\sum_{k=1}^L \alpha_k p_k(x | \theta_k)}$$

Posterior probability

GMM based segmentation

- The M Step - parameter updates

$$\text{weight } \alpha_l^{(m+1)} = \frac{1}{n} \sum_{j=1}^n P(l | x_j, \theta_l^{(m)})$$

$$\text{Mean } \mu_l^{(m+1)} = \frac{\sum_{j=1}^n x_j P(l | x_j, \theta_l^{(m)})}{\sum_{j=1}^n P(l | x_j, \theta_l^{(m)})}$$

$$\text{Covariance } \Sigma_l^{(m+1)} = \frac{\sum_{j=1}^n P(l | x_j, \theta_l^{(m)}) \{ (x_j - \mu_l^{(m)})(x_j - \mu_l^{(m)})^T \}}{\sum_{j=1}^n P(l | x_j, \theta_l^{(m)})}$$

Summary of segmentation

We have seen simple to complex methods

- Pixel and region based approaches
- Thresholding, clustering
- Iterative and non-iterative methods

All of the above use a bottom-up approach

Recent approaches are based on

- Partial differential equations (snakes, active contours, level sets..)
- Graph partitioning (graph cuts..)

These are covered in **CV, MIP** courses!