

# Statistical Methods in Artificial Intelligence

## CSE471 - Monsoon 2016 : Lecture 08




Avinash Sharma  
CVIT, IIIT Hyderabad

# Lecture Plan

- Recap
- Backpropagation as Feature Mapping
- Practical Aspects of Backpropagation
- Additional Networks
  - Deep Learning & Convolution Networks (ConvNet)
  - Recurrent Networks
  - RBF Networks


# Backpropagation in NN



# Backpropagation in Neural Networks

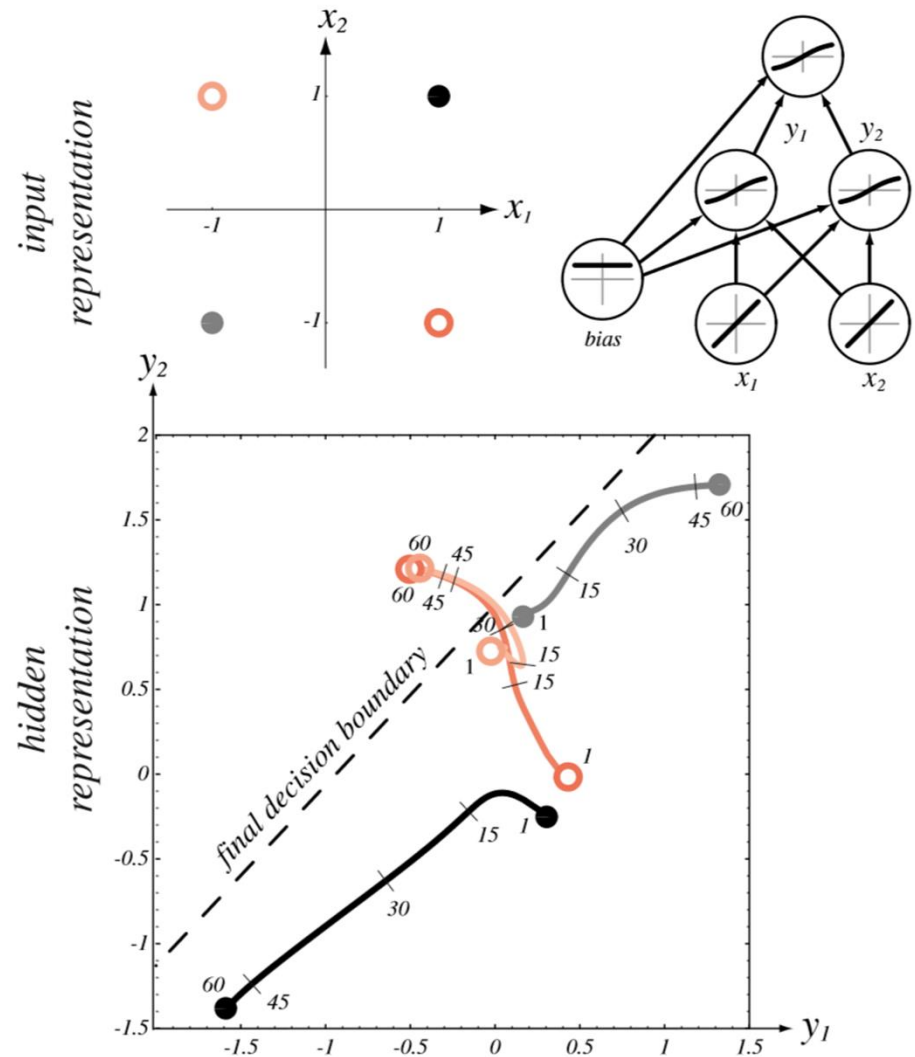
- $J(\mathbf{w}) = \frac{1}{2} \sum_{k=1}^c (t_k - z_k)^2 = \frac{1}{2} \|\mathbf{t} - \mathbf{z}\|^2$
- $\Delta \mathbf{w} = -\eta \frac{\partial J}{\partial \mathbf{w}}, \Delta w_{pq} = -\eta \frac{\partial J}{\partial w_{pq}}$
- $\Delta w_{kj} = \eta \delta_k y_j = \eta (t_k - z_k) f'(net_k) y_j$
- $\Delta w_{ji} = \eta \delta_j x_i = \eta \left[ \sum_{k=1}^c w_{kj} \delta_k \right] f'(net_j) x_i$

# Modelling the Non-linearity




# Backpropagation as Feature Mapping

- Output of hidden layers turns out to be linearly separable.
- Input-hidden layer achieves non-linear transform.
- Hidden-output layer feed forward only achieves a linear classification.



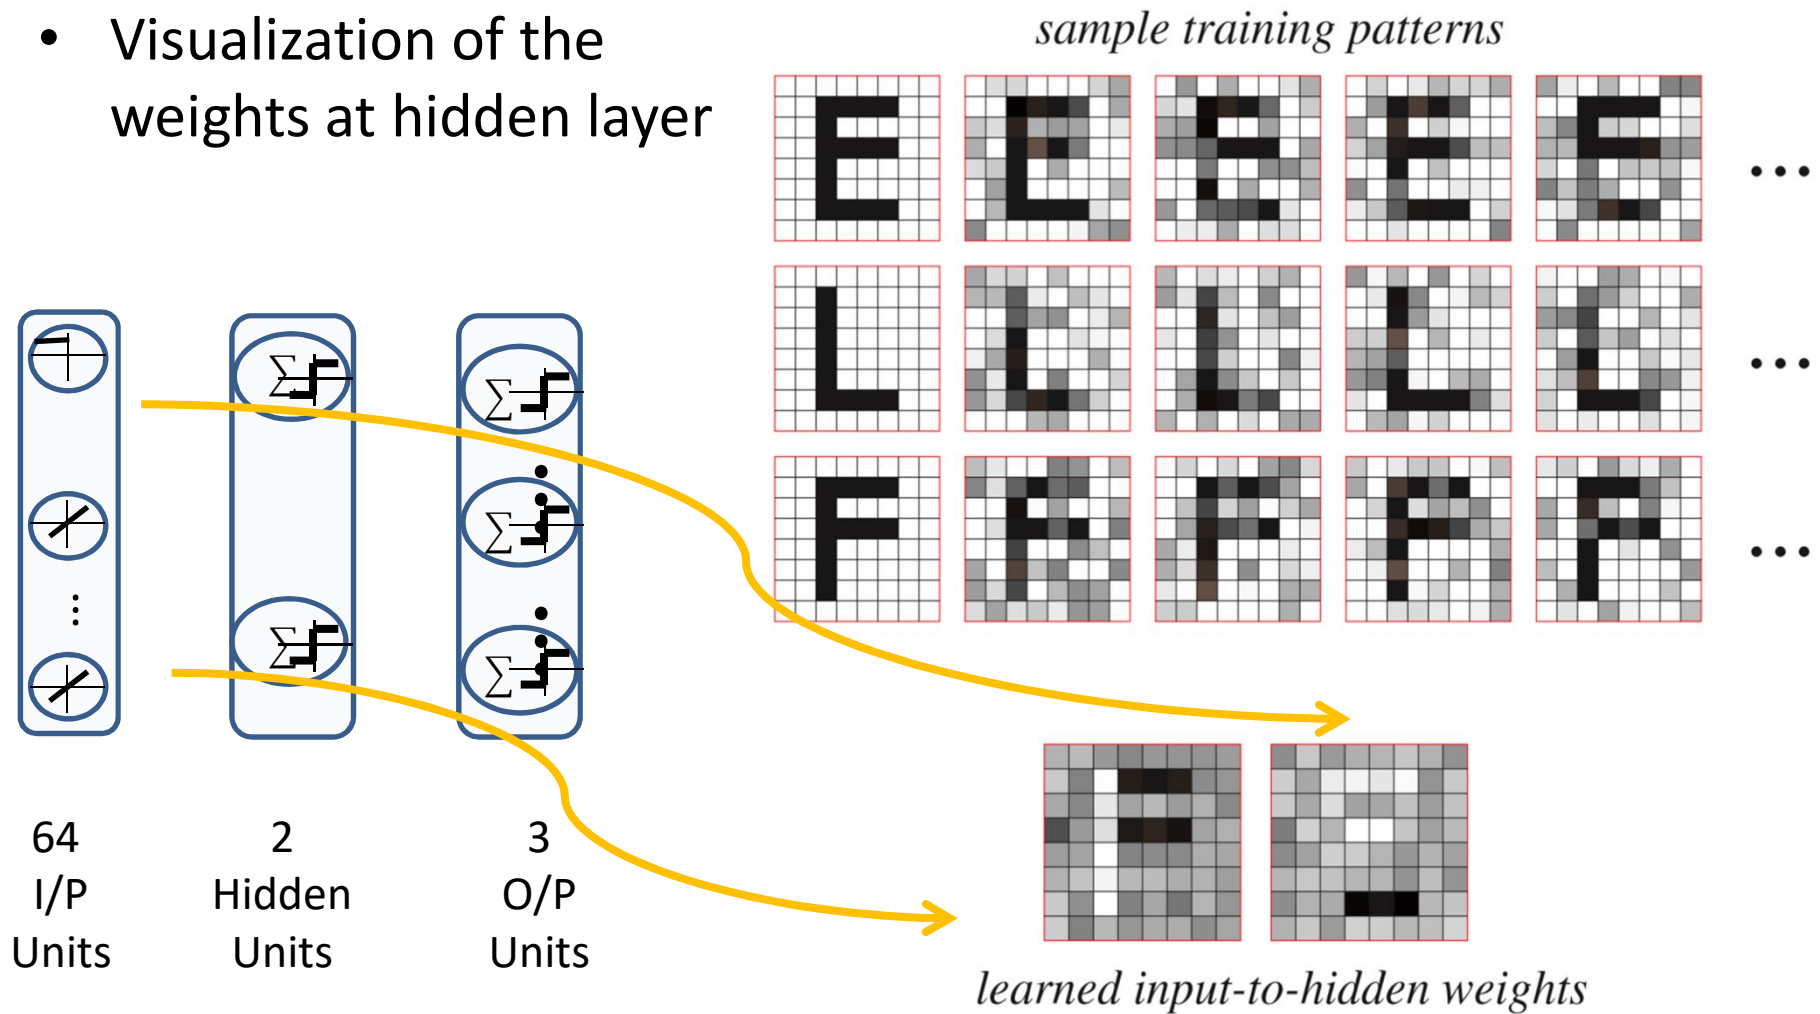
# Backpropagation as Feature Mapping

- Output of hidden layers turns out to be linearly separable.
- Input-hidden layer achieves non-linear transform.
- Hidden-output layer feed forward only achieves a linear classification.
- Therefore, adding more hidden units might improve the performance



# Backpropagation as Feature Mapping

- Visualization of the weights at hidden layer






# Practical Aspects of Backpropagation

- Activation Function
  - $f(\cdot)$  should be ***Non-linear***
  - $f(\cdot)$  should ***Saturate***
  - $f(\cdot)$  should be ***Continuous & Smooth***
  - $f'(\cdot)$  should be ***Defined***
  - $f(\cdot)$  can have ***Monotonicity***
  - $f(\cdot)$  can be ***linear for small values of net***

# Practical Aspects of Backpropagation

- **Scaling of Input**

- $X = [x_1, \dots, x_m]^T_{m \times d}$
- $\tilde{X} = X - \text{mean}(X)$  (Centering of Data)
- $X_{Norm} = \tilde{X} / \sigma$  (Column-wise division by Standard Deviation of each dimension)



# Practical Aspects of Backpropagation

- **Scaling of Input**

- $X = [\mathbf{x}_1, \dots, \mathbf{x}_m]^T_{m \times d}$
- $\tilde{X} = X - \text{mean}(X)$  (Centering of Data)
- $X_{Norm} = \tilde{X} / \sigma$  (Column-wise division by Standard Deviation of each dimension)

- **Target Values**

- Use +1 and −1 or any real value in this range as output
- Related to saturation value of the activation function

- **Training with Noise**

- Add random noise to original training samples for generating more training samples

# Practical Aspects of Backpropagation

- **Manufacturing Data**

- Add translation and rotation transforms to original training data to generate more rich training data samples

- **Number of Hidden Units**

- Too few leads to high test error due to lack of expressibility
- Too many leads to overfitting to training data
- Choose such that total number of weights =  $m/10$ .

- **Weight Initialization**



- Do not initialize with zero weights
- Use both random positive & negative weights as data is standardized
- $-1/\sqrt{d} < w_{ji} < +1/\sqrt{d}$  and  $-1/\sqrt{nH} < w_{kj} < +1/\sqrt{nH}$

# Practical Aspects of Backpropagation

- **Learning Rates**

- For quadratic error criterion function  $J$  :

$$\eta_{opt} = \left( \frac{\partial^2 J}{\partial \mathbf{w}^2} \right)^{-1}$$



# Practical Aspects of Backpropagation

- **Learning Rates**


- For quadratic error criterion function  $J$  :

$$\eta_{opt} = \left( \frac{\partial^2 J}{\partial \mathbf{w}^2} \right)^{-1}$$

- **Momentum**

- Continue with inertia in each weight update
- $\mathbf{w}(m + 1) = \mathbf{w}(m) + (1 - \alpha)\Delta\mathbf{w}(m) + \alpha\Delta\mathbf{w}(m - 1)$

# Practical Aspects of Backpropagation



# Practical Aspects of Backpropagation

- **Learning Rates**

- For quadratic error criterion function  $J$  :

$$\eta_{opt} = \left( \frac{\partial^2 J}{\partial \mathbf{w}^2} \right)^{-1}$$

- **Momentum**

- Continue with inertia in each weight update
- $\mathbf{w}(m + 1) = \mathbf{w}(m) + (1 - \alpha)\Delta\mathbf{w}(m) + \alpha\Delta\mathbf{w}(m - 1)$


- **Weight Decay**

- $w_{new} = w_{old}(1 - \epsilon)$
- Weights that do not affect the error function will eventually become zero.



# Practical Aspects of Backpropagation

- Hints
  - Add ancillary units to output only for training phase.
  - In testing phase remove these extra units and related weights.



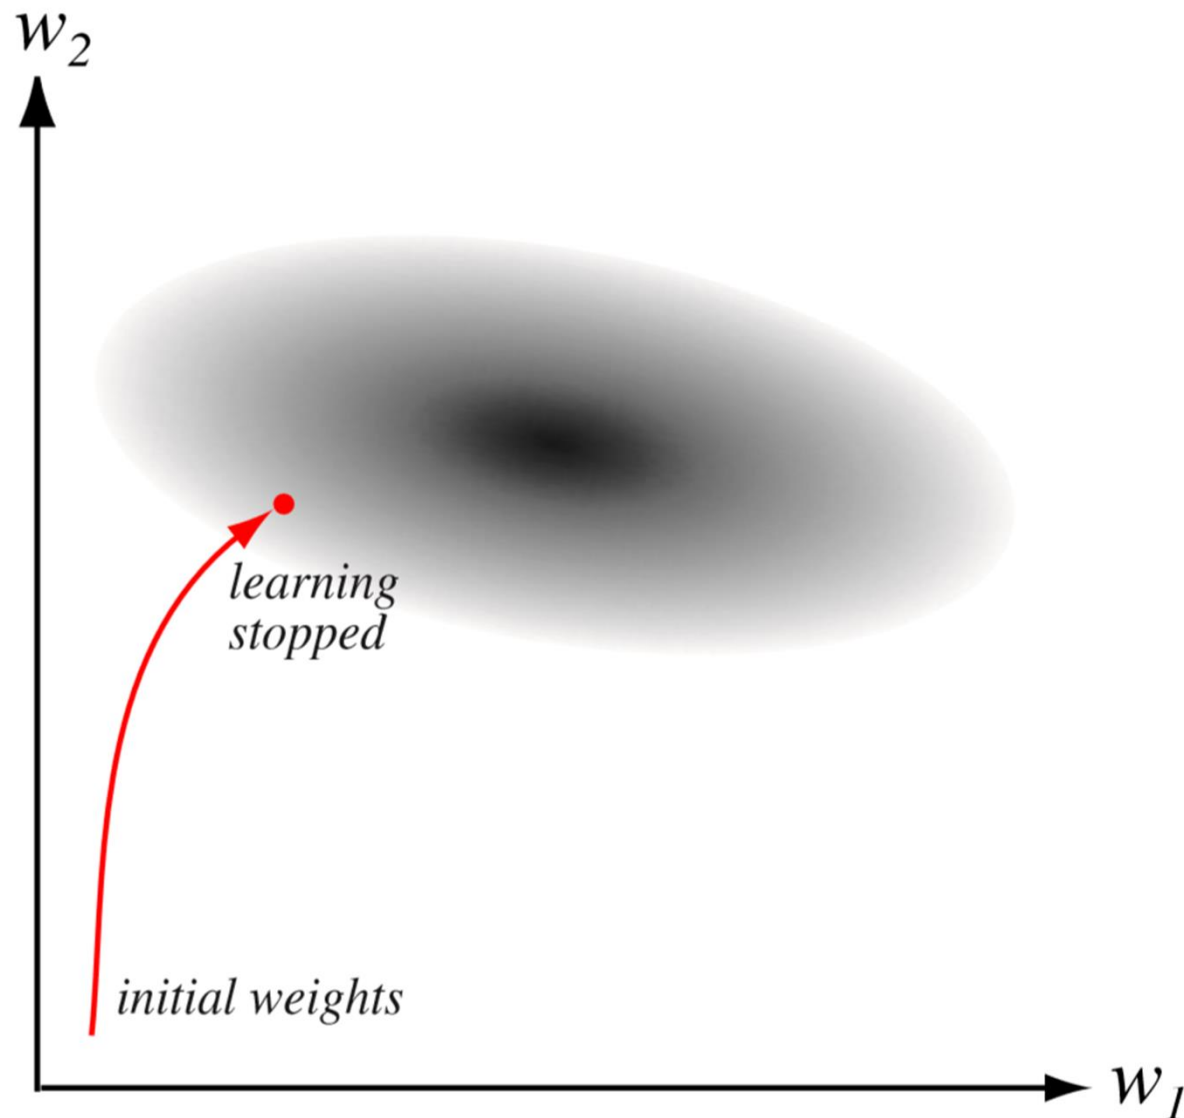
# Practical Aspects of Backpropagation

- Hints
  - Add ancillary units to output only for training phase.
  - In testing phase remove these extra units and related weights
- Online/Batch/Stochastic Training
  - Stochastic training is mostly preferred over batch
  - Online is rarely used for tasks where storing all data samples is prohibitive

# Practical Aspects of Backpropagation

- Hints
  - Add ancillary units to output only for training phase.
  - In testing phase remove these extra units and related weights
- Online/Batch/Stochastic Training
  - Stochastic training is mostly preferred over batch
  - Online is rarely used for tasks where storing all data samples is prohibitive
- Stopped Training
  - Excessive training can lead to poor generalization
  - Stopping training before low error reached is good to avoid overfitting

# Practical Aspects of Backpropagation




# Practical Aspects of Backpropagation

- Number of Hidden Layers
  - Depends on complexity of classification problem
  - Unnecessary layers can cause minimization to caught into local minima
- Criterion Function
  - Entropy based information theoretic error functions
  - Minkowski error function (generalization of sum of squared error)


# Convolution Networks

- More suitable for image data.




# Convolution Networks

- <http://cs231n.github.io/convolutional-networks/>



# Convolution Networks


- <http://cs231n.github.io/convolutional-networks/>





# Convolution Networks

- Live Demo at <http://cs231n.stanford.edu/>




# Convolution Networks

- **INPUT** [32x32x3] will hold the raw pixel values of the image,.
- **CONV** [32x32x12] layer will compute the output of neurons that are connected to local regions in the input, each computing a dot product between their weights and the region they are connected to in the input volume.
- **RELU** layer will apply an elementwise activation function
- **POOL** [16x16x12] layer will perform a downsampling operation along the spatial dimensions (width, height).
- **FC** (i.e. fully-connected) layer will compute the class scores, resulting in volume of size [1x1x10], where each of the 10 numbers correspond to a class score.

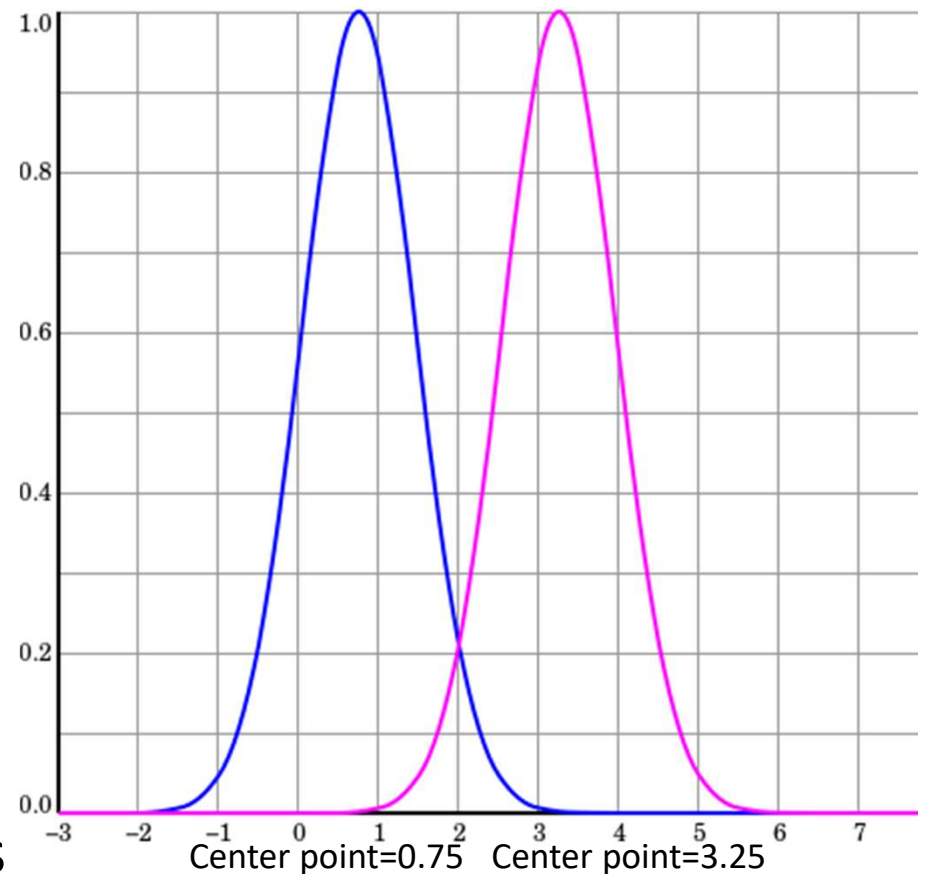
# Recurrent Networks

- Useful for time-dependent signals with short periodic structures

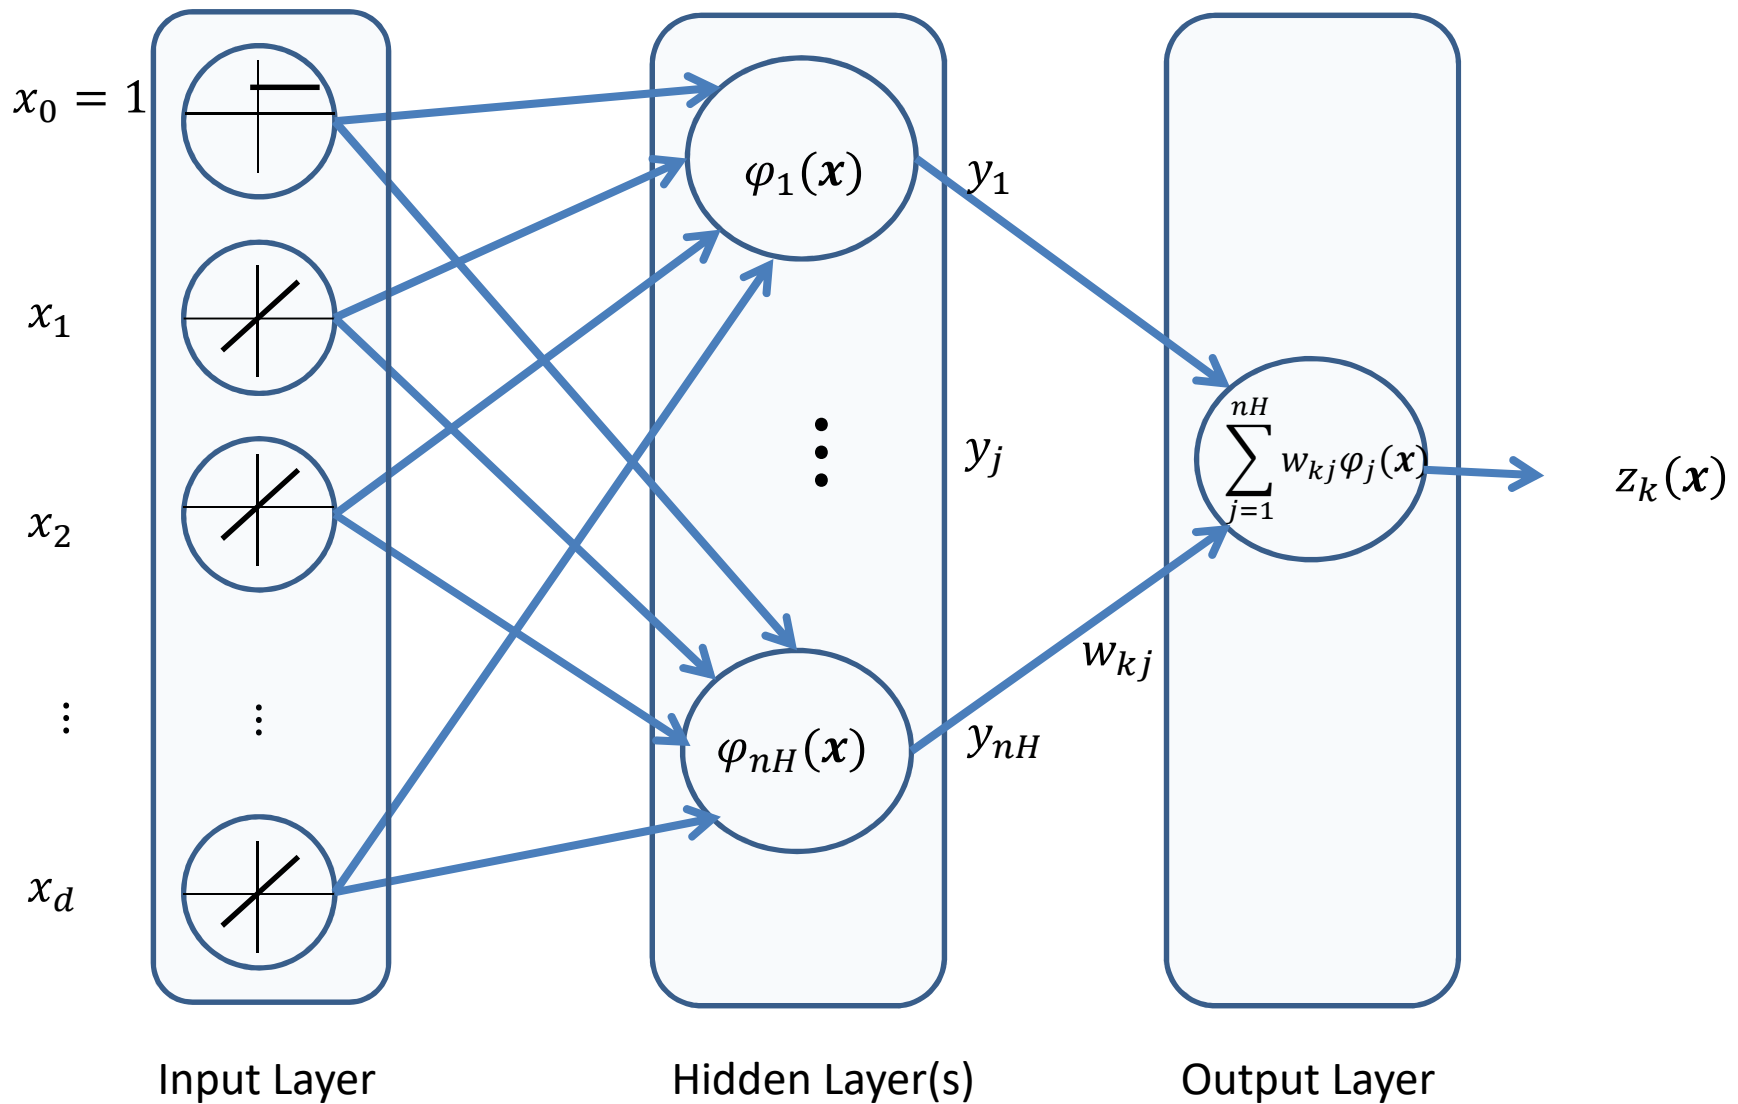


# RBF Networks

- Radial functions are a special class of function
- Their characteristic feature is that their response decreases (or increases) monotonically with distance from a center point
- The center, the distance scale, and the precise shape of the radial function are parameters of the model.



# RBF Networks



# RBF Networks

- $z_k(\mathbf{x}) = \sum_{j=1}^{nH} w_{kj} \varphi_j(\mathbf{x})$
- Let  $\Phi_{m \times nH} = \begin{bmatrix} \varphi_1(\mathbf{x}_1) & \dots & \varphi_{nH}(\mathbf{x}_1) \\ \vdots & \dots & \vdots \\ \varphi_1(\mathbf{x}_m) & \dots & \varphi_{nH}(\mathbf{x}_m) \end{bmatrix}$
- Let  $T_{m \times 1} = [t_1 \quad \dots \quad t_m]^T$
- $J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^m \|z_i - t_i\|^2$
- $\Phi^T \Phi W^T = \Phi^T T$
- $W^T = (\Phi^T \Phi)^{-1} \Phi^T T = \Phi^\dagger T$
- BP can be used if  $z_k(\mathbf{x}) = f\left(\sum_{j=1}^{nH} w_{kj} \varphi_j(\mathbf{x})\right)$  for any non-linear  $f$